

# Graph Neural Networks

---

Natalia Dubljevic  
ENEL 645 Co-instructor  
Biomedical Engineering  
Schulich School of Engineering

W2024

# Outline

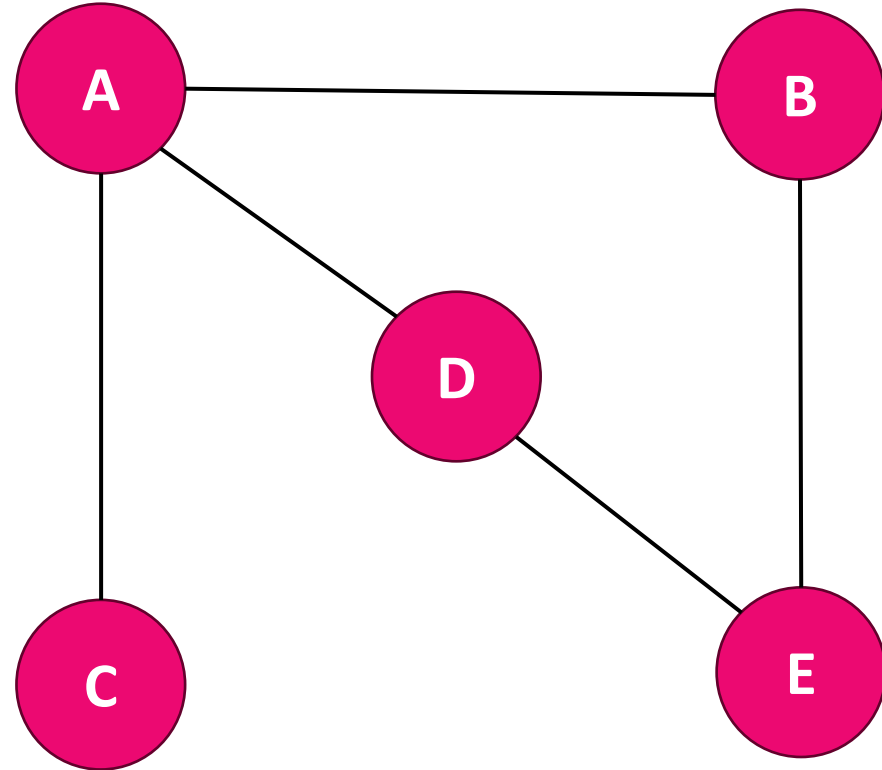
- Learning Goals
- Graph neural networks
  - What is graph-structured data
  - What sorts of problems can we solve using graph neural networks
  - How are graph neural networks designed
- Summary

# Learning Goals

- Understand graph structured data and its differences from data we've seen so far
- What are some applications of GNNs
- Some of the core concepts of GNN implementation such as
  - Message passing
  - K-hop neighbourhood

# What is a graph anyways?

- Graphs have **nodes** and **edges**
- Edges represent the relationship between any two nodes



# Graph structured data

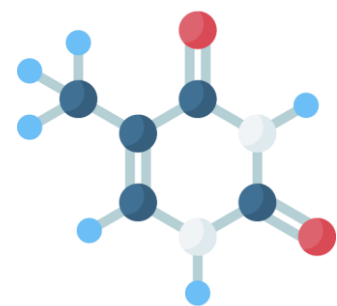
Graphs can be used to represent data such as:

## Social networks



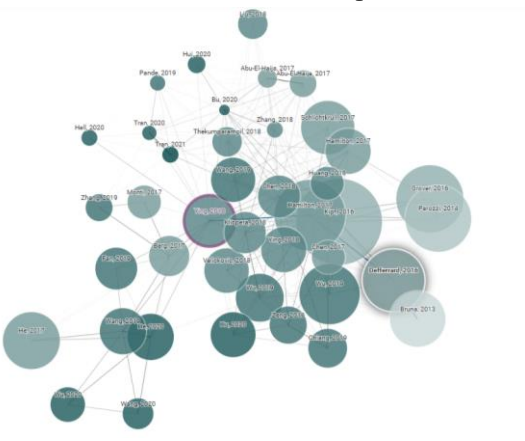
[University of Kentucky]

## Molecules



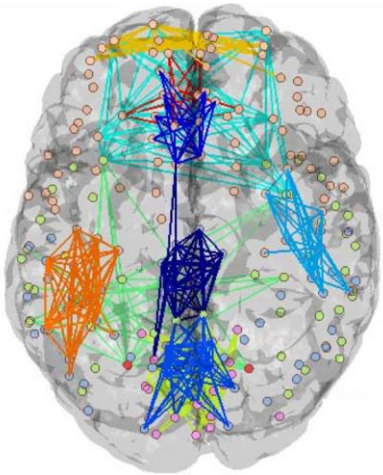
[flaticon.com]

## Citation Maps



[connectedpapers.com]

## Brain Connectivity



[spectrumnews.org]

## Transit

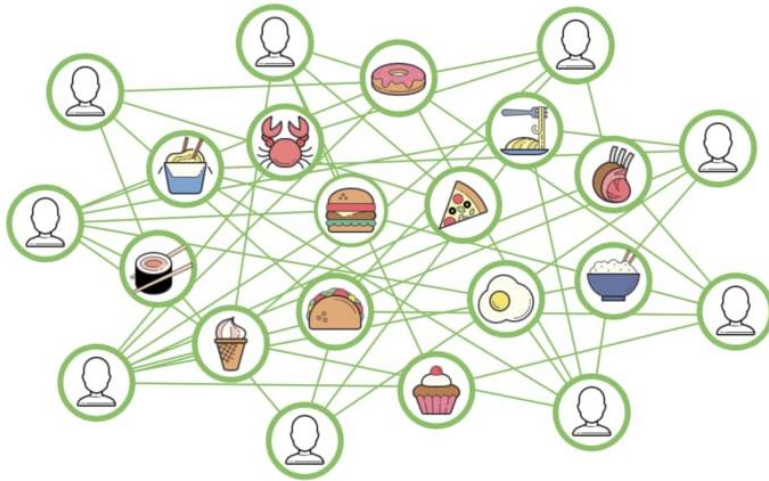


[translink.ca]

# Graph structured data

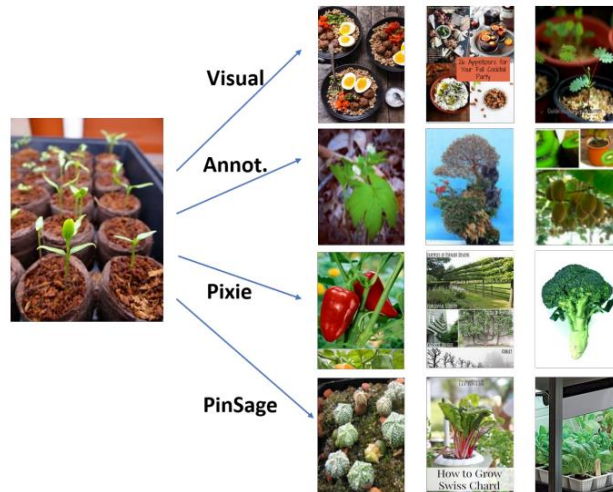
GNNs aren't just for niche applications-- they're trendy in industry too!

## Uber Eats



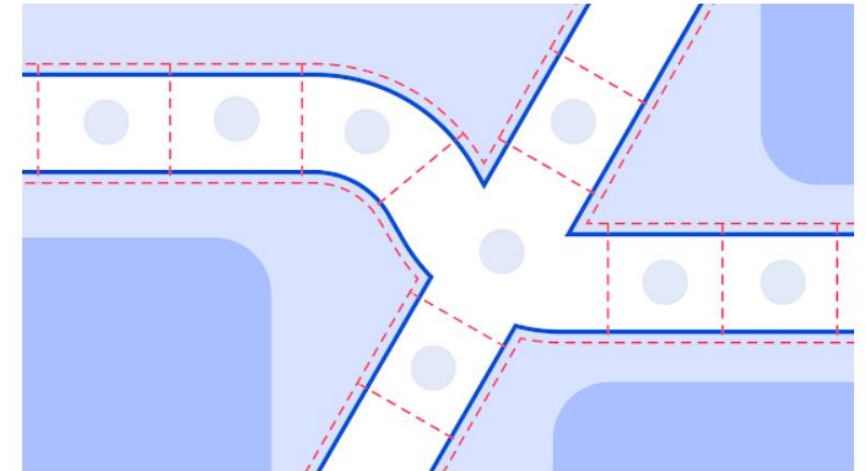
[uber.com]

## Pinterest



[Ying et al. 2018]

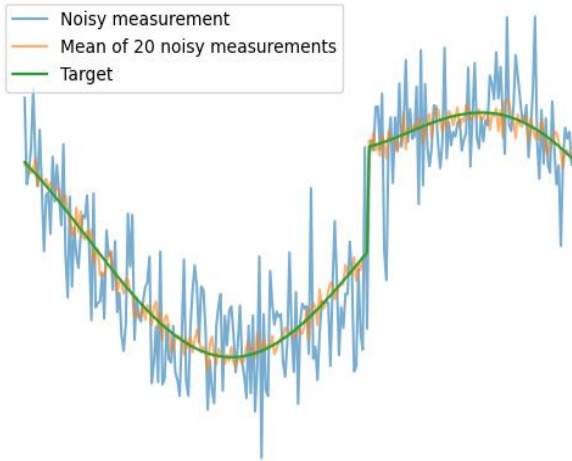
## Google maps



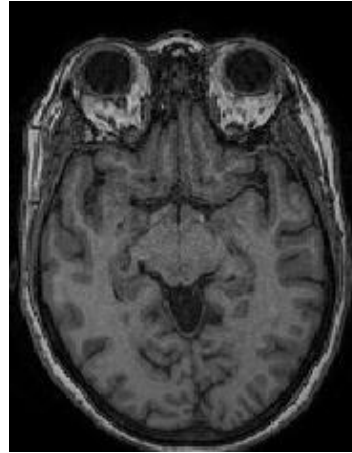
[deepmind.google.com]

# What we've seen so far

## Signals



## Images



## Tabular

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife

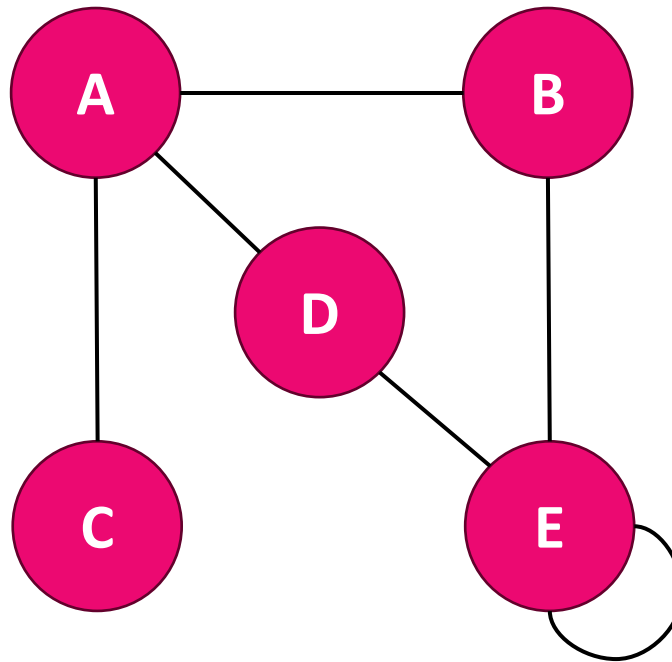
## Text

Love looks not with the eyes, but  
with the mind; And therefore is  
wing'd Cupid painted blind. Nor  
hath love's mind of any judgment  
taste; Wings and no eyes figure  
unheedy haste: And therefore is  
love said to be a child, Because in  
choice he is so oft beguil'd.

- All of these can be represented as graphs!
- Graphs are one of the most general data structures

# A bit more on graphs

- We can represent graphs as an adjacency matrix

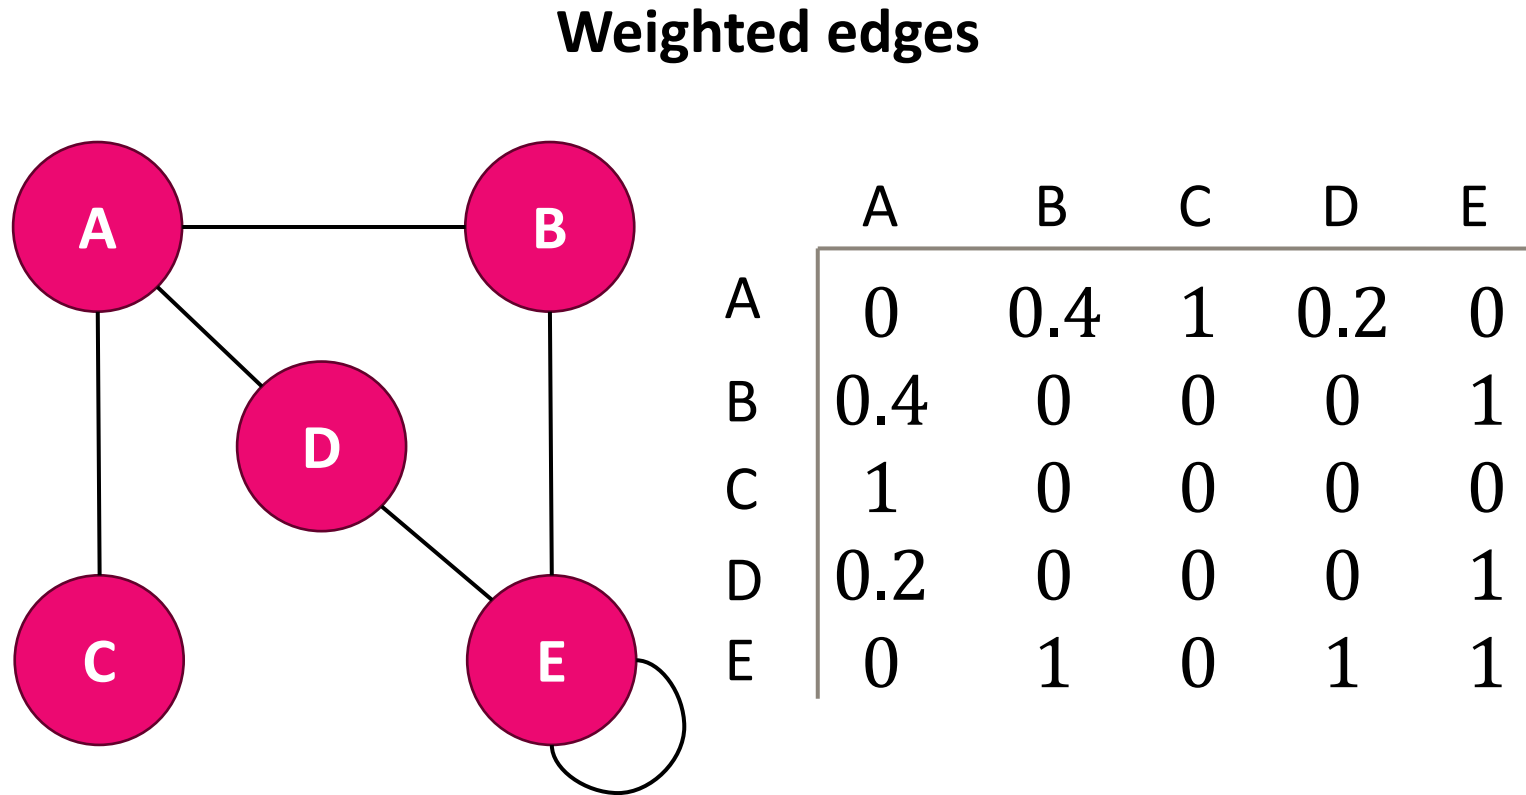


	A	B	C	D	E
A	0	1	1	1	0
B	1	0	0	0	1
C	1	0	0	0	0
D	1	0	0	0	1
E	0	1	0	1	1



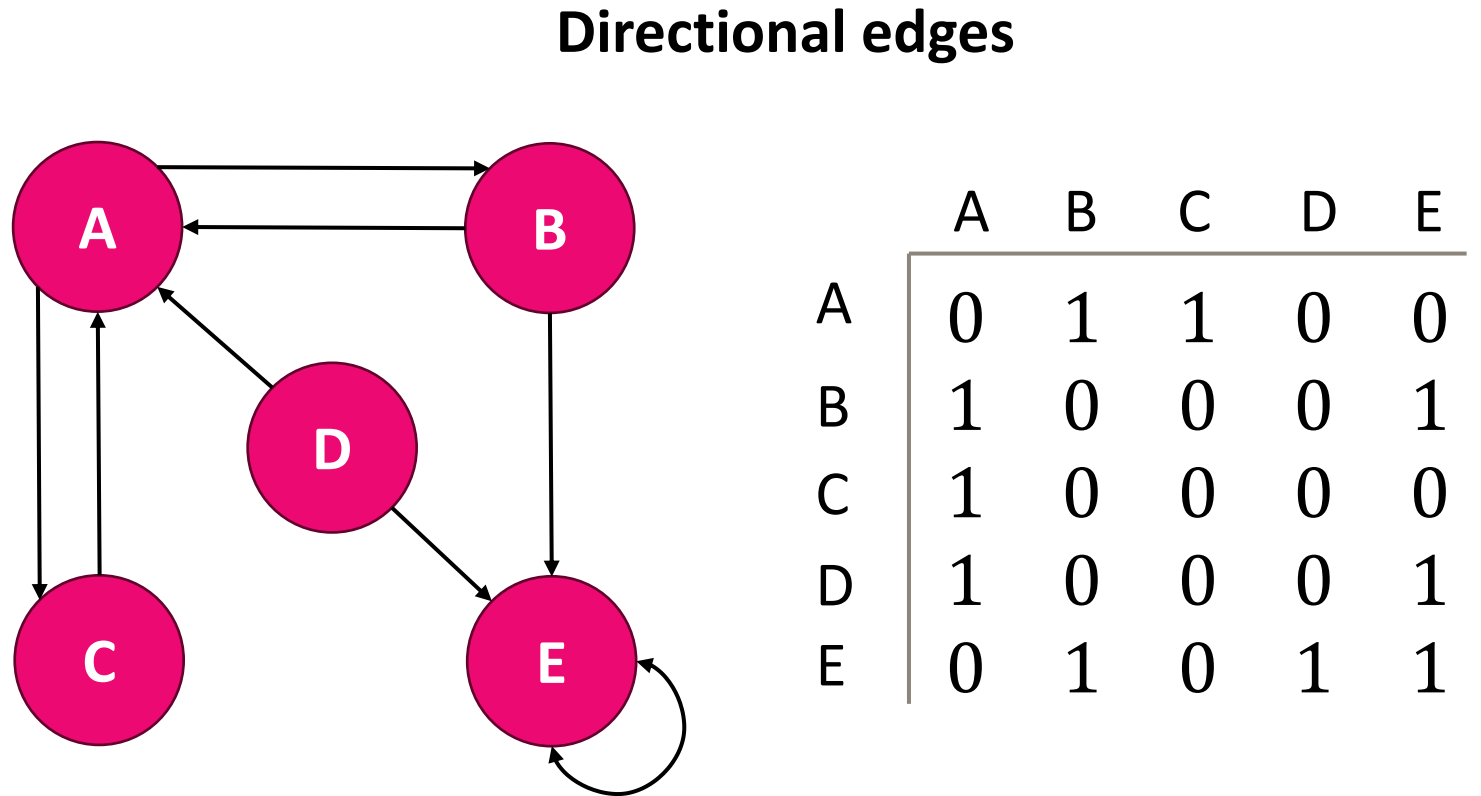
# A bit more on graphs

- We can represent graphs as an adjacency matrix
- Edges can be weighted and/or directional



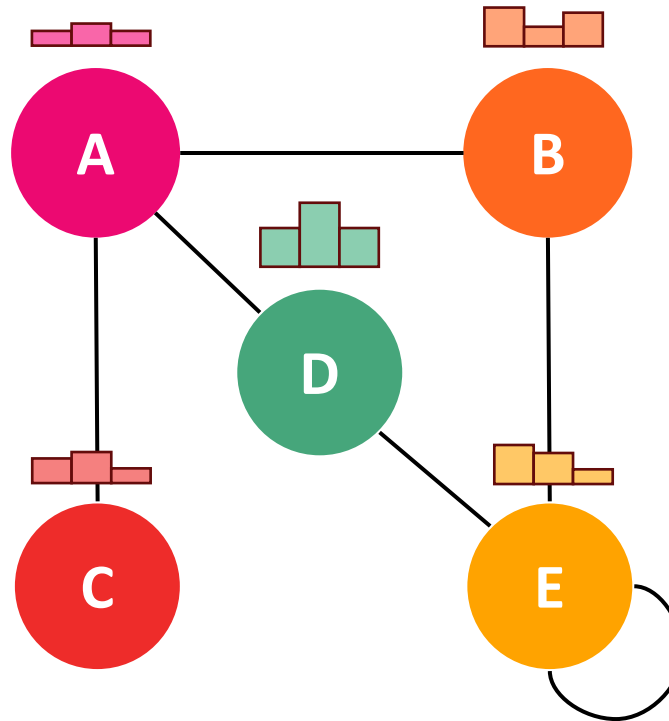
# A bit more on graphs

- We can represent graphs as an adjacency matrix
- Edges can be weighted and/or directional



# A bit more on graphs

- We can represent graphs as an adjacency matrix
- Edges can be weighted and/or directional
- Can have features related to nodes, edges, or graph as a whole
  - Node features most common



	A	B	C	D	E
A	0	1	1	1	0
B	1	0	0	0	1
C	1	0	0	0	0
D	1	0	0	0	1
E	0	1	0	1	1

# Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	A	B	C	D	E
A	0	1	1	1	0
B	1	0	0	0	1
C	1	0	0	0	0
D	1	0	0	0	1
E	0	1	0	1	1

# Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

# Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$(0, 1)$

# Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$\left[ \begin{array}{l} (0, 1) \\ (0, 2) \end{array} \right]$

# Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$$\begin{pmatrix} (0, 1) \\ (0, 2) \\ (0, 3) \end{pmatrix}$$



# Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$$\begin{pmatrix} (0, 1) \\ (0, 2) \\ (0, 3) \\ (1, 0) \end{pmatrix}$$

# Working with sparse tensors

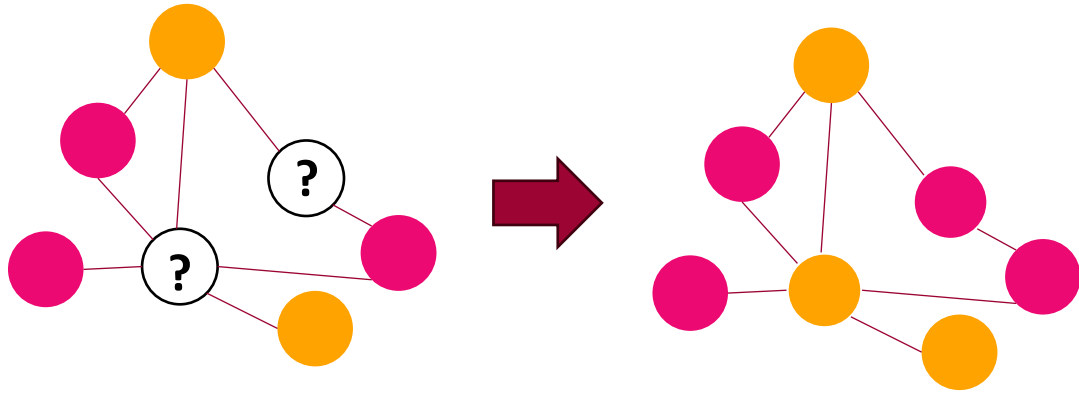
- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

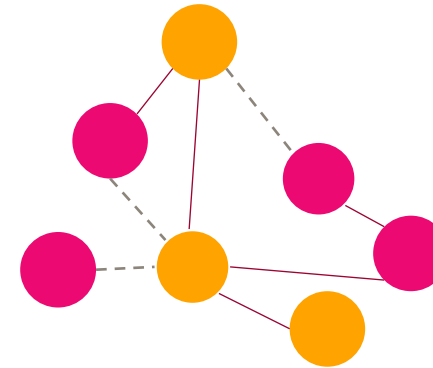
$$\begin{pmatrix} (0, 1) \\ (0, 2) \\ (0, 3) \\ (1, 0) \\ (1, 4) \\ \vdots \end{pmatrix}$$

# Tasks

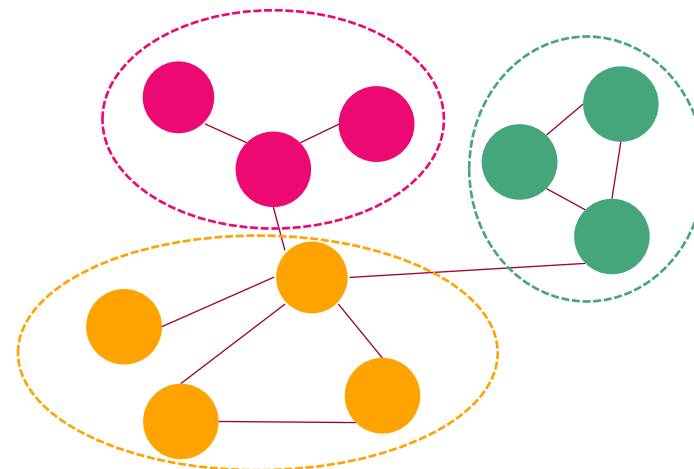
## Node Classification



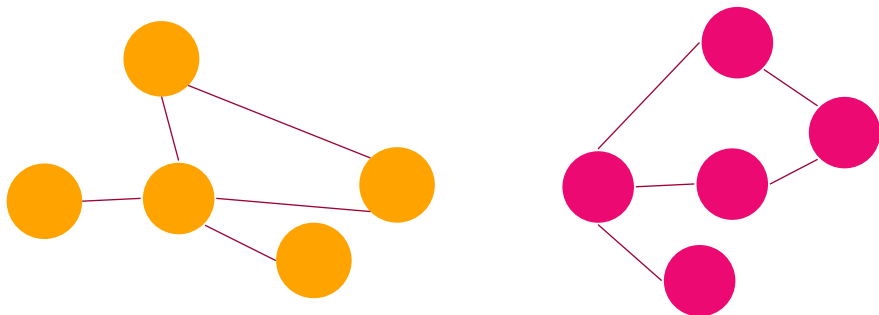
## Edge prediction



## Clustering

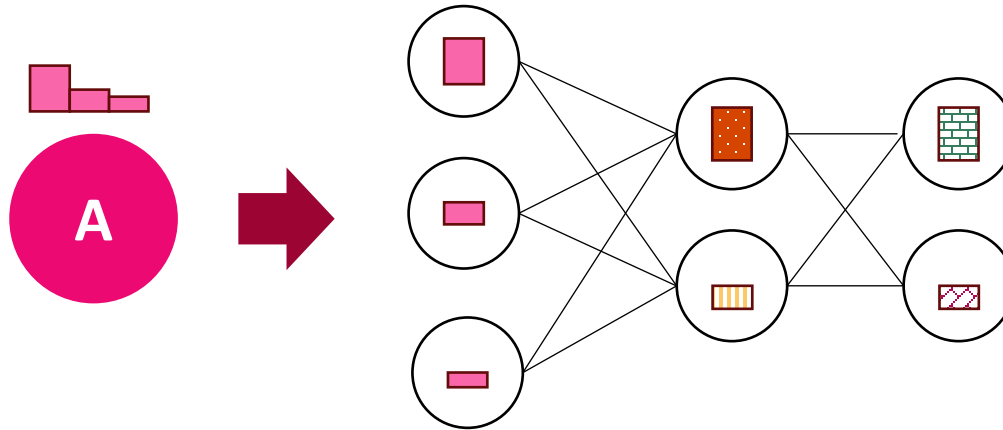


## Graph Classification



# Prelude to message passing

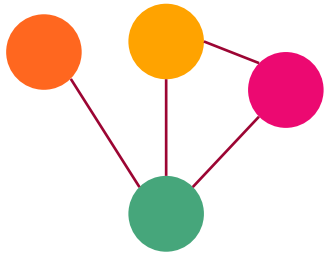
- How do we create models that can accommodate graph-structured data?
- Naïve approach: Feed a node and its features into an FCNN



- **Question:** Are there any issues with this approach?

# Prelude to message passing

- Unlike standard supervised learning, we do not have i.i.d. points!
  - Our nodes are connected to each other in unique ways

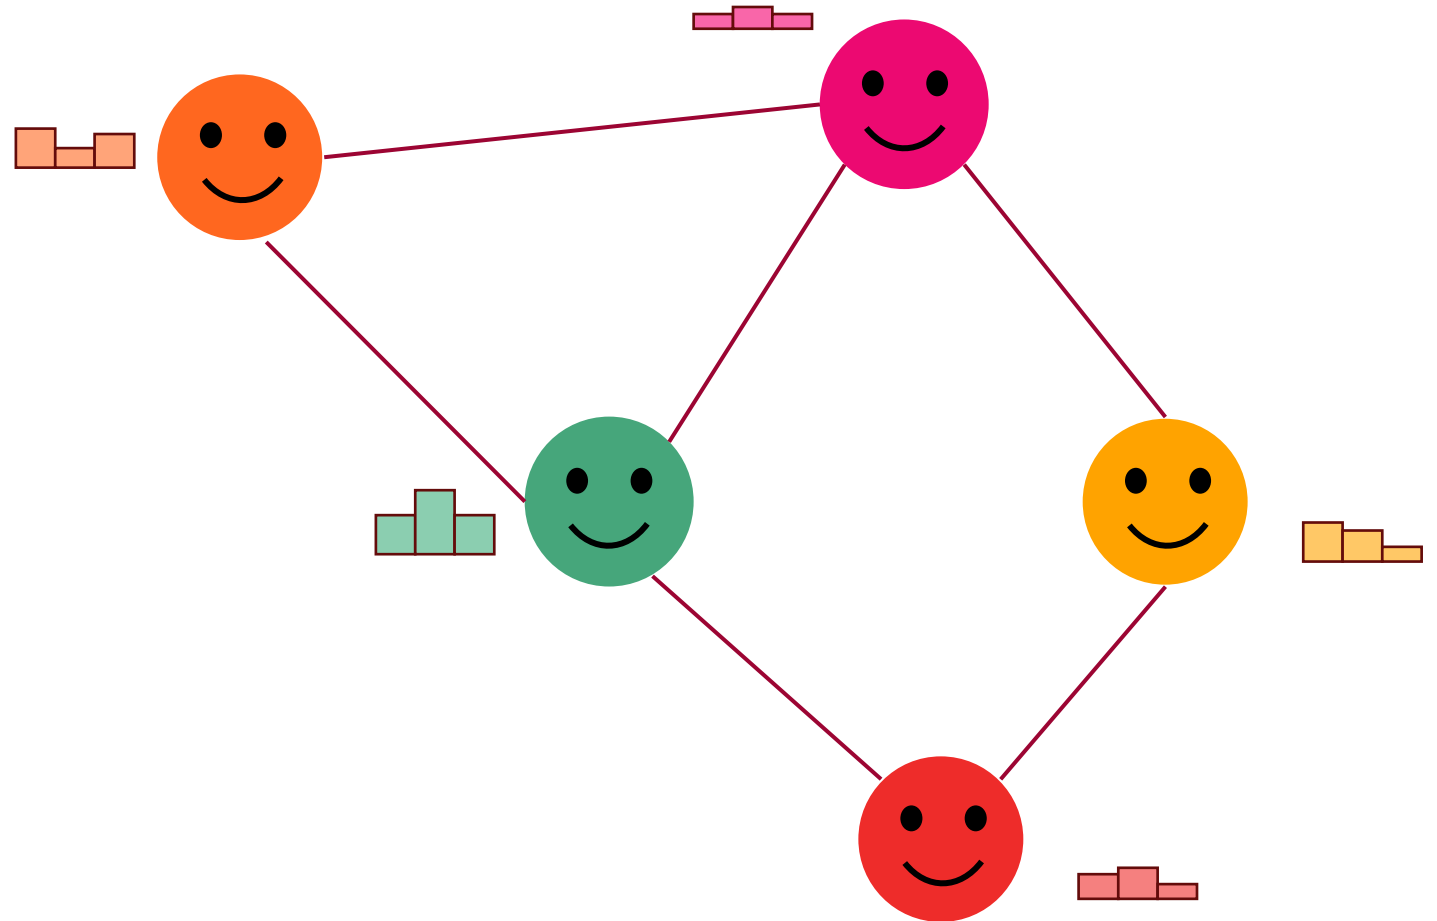


is not the  
same as



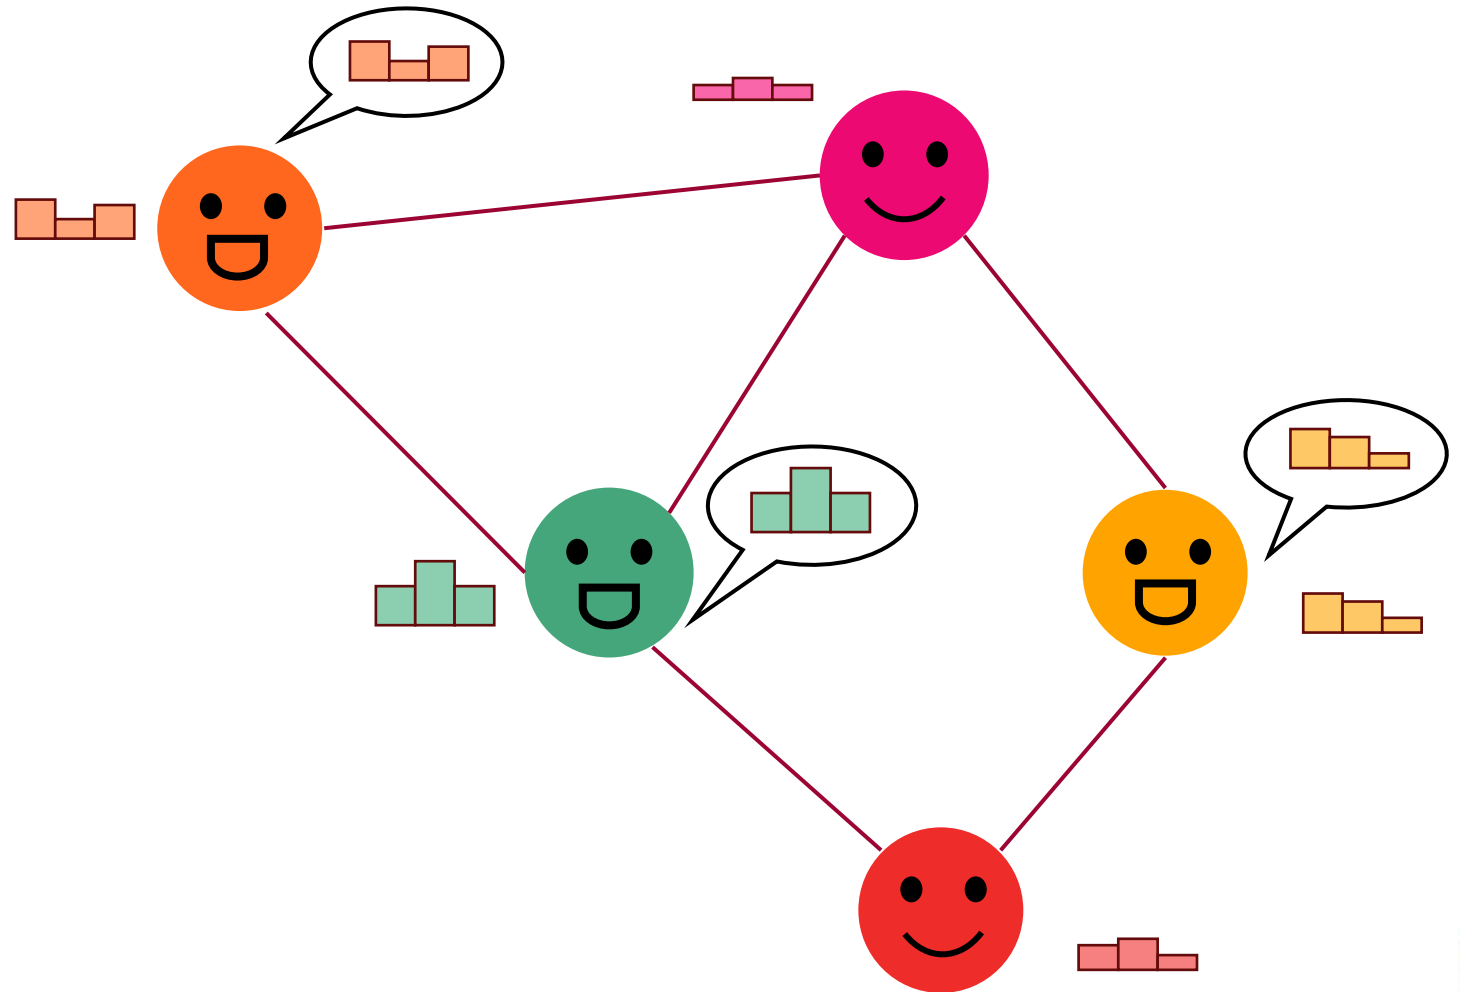
- We want to develop models that can exploit the relationships between the nodes
- We'll introduce '**message passing**' which allows for information to be propagated effectively throughout a neural network

# Message passing



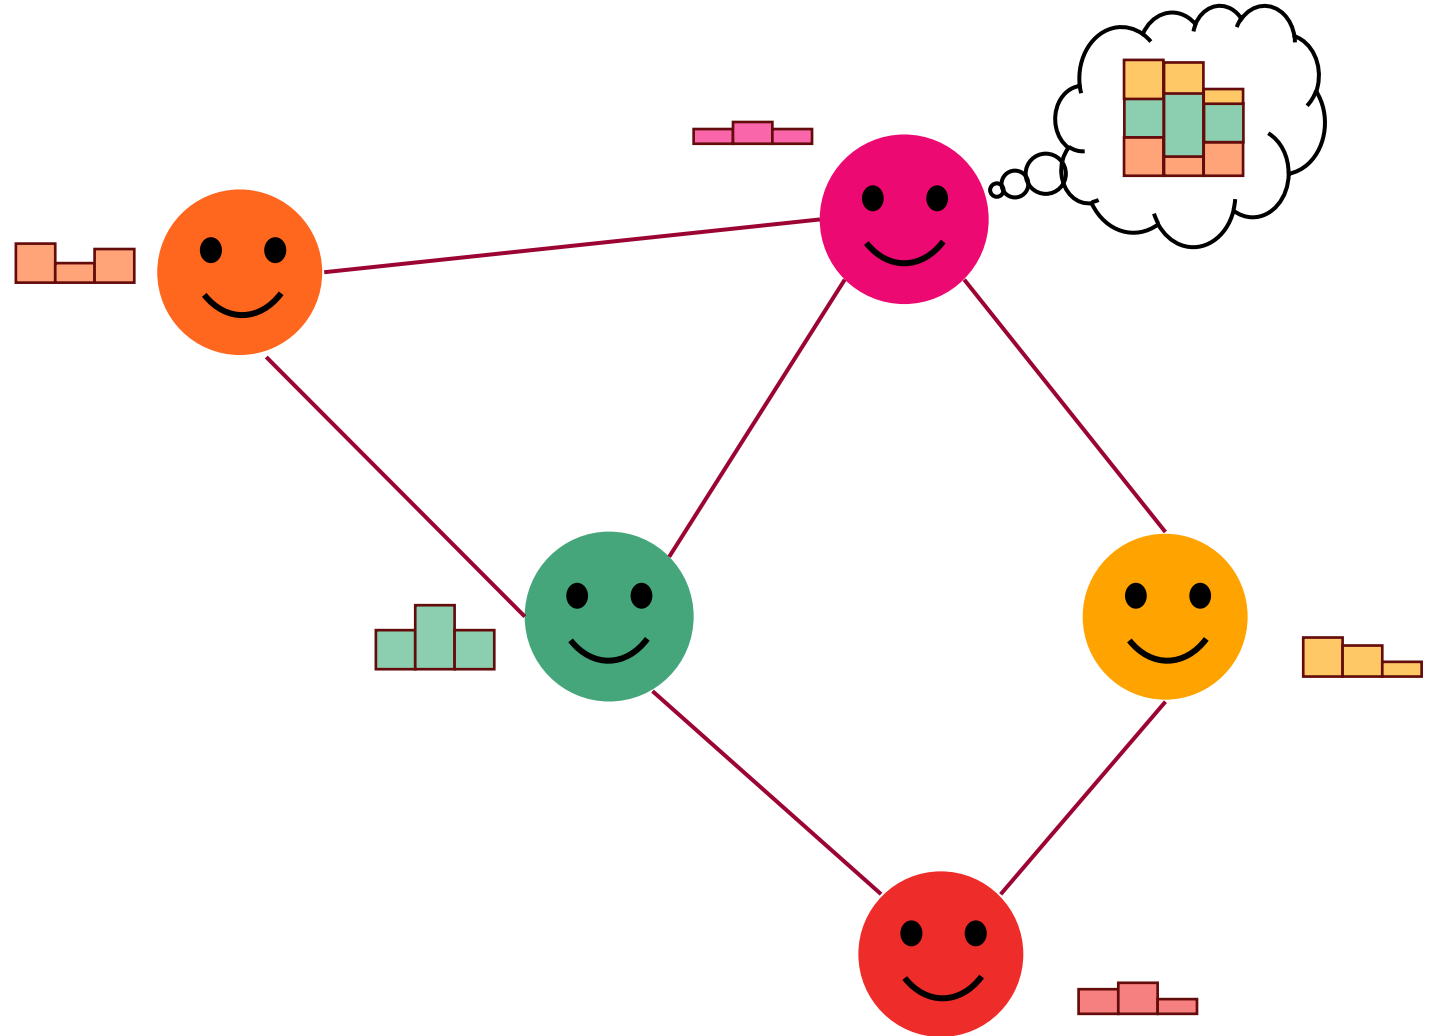
# Message passing

1. **Message:** Nodes send messages to their neighbors



# Message passing

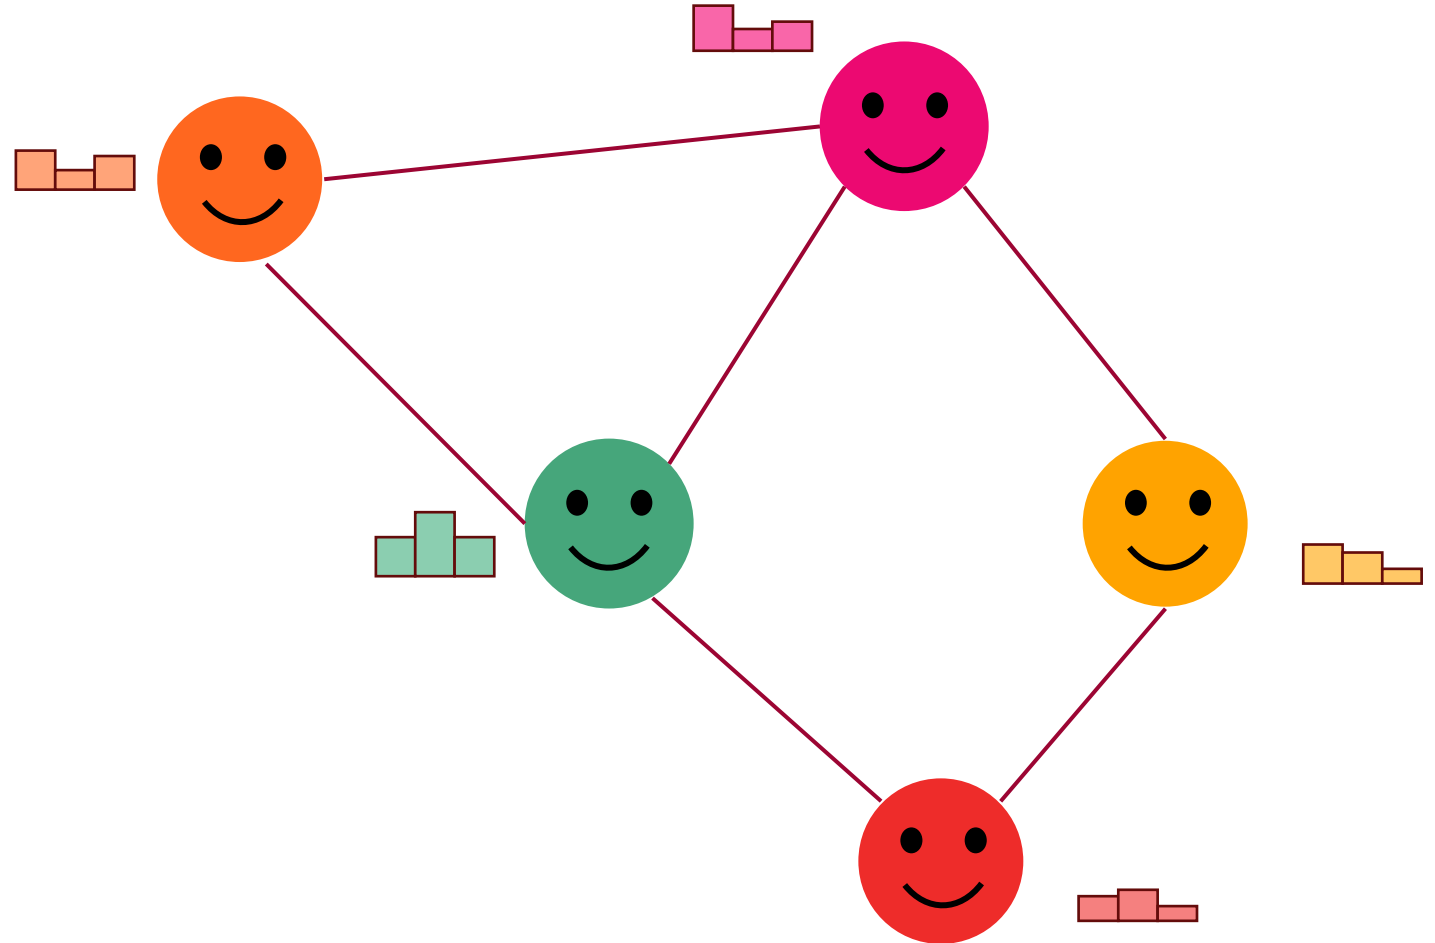
1. **Message:** Nodes send messages to their neighbors
2. **Aggregate:** Messages are aggregated in a permutation-invariant way





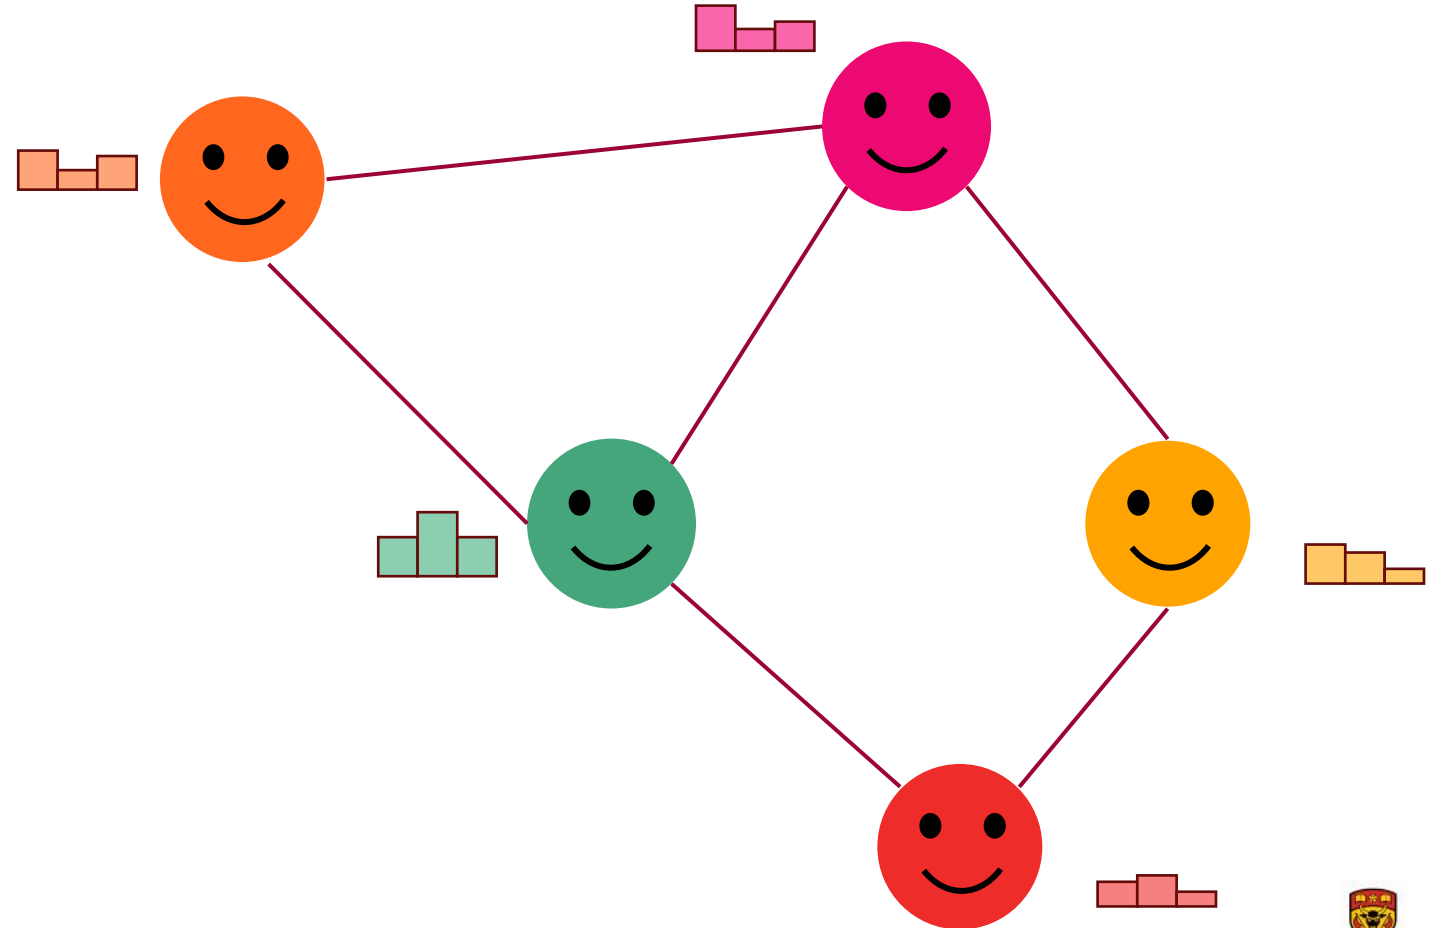
# Message passing

1. **Message:** Nodes send messages to their neighbors
2. **Aggregate:** Messages are aggregated in a permutation-invariant way
3. **Update:** Node embeddings are updated



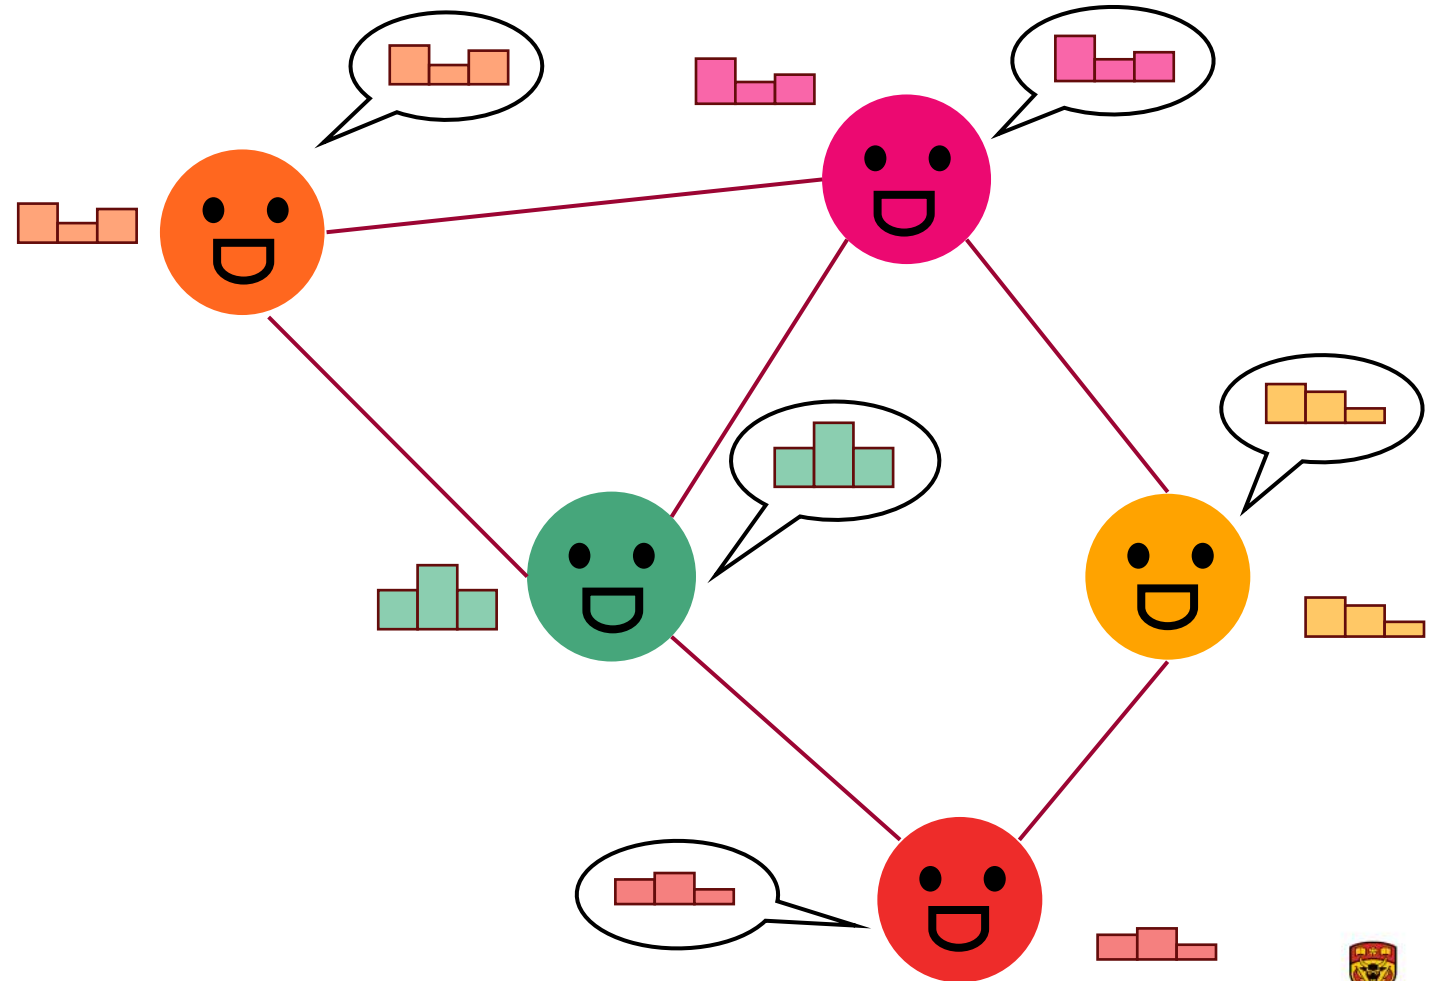
# Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!



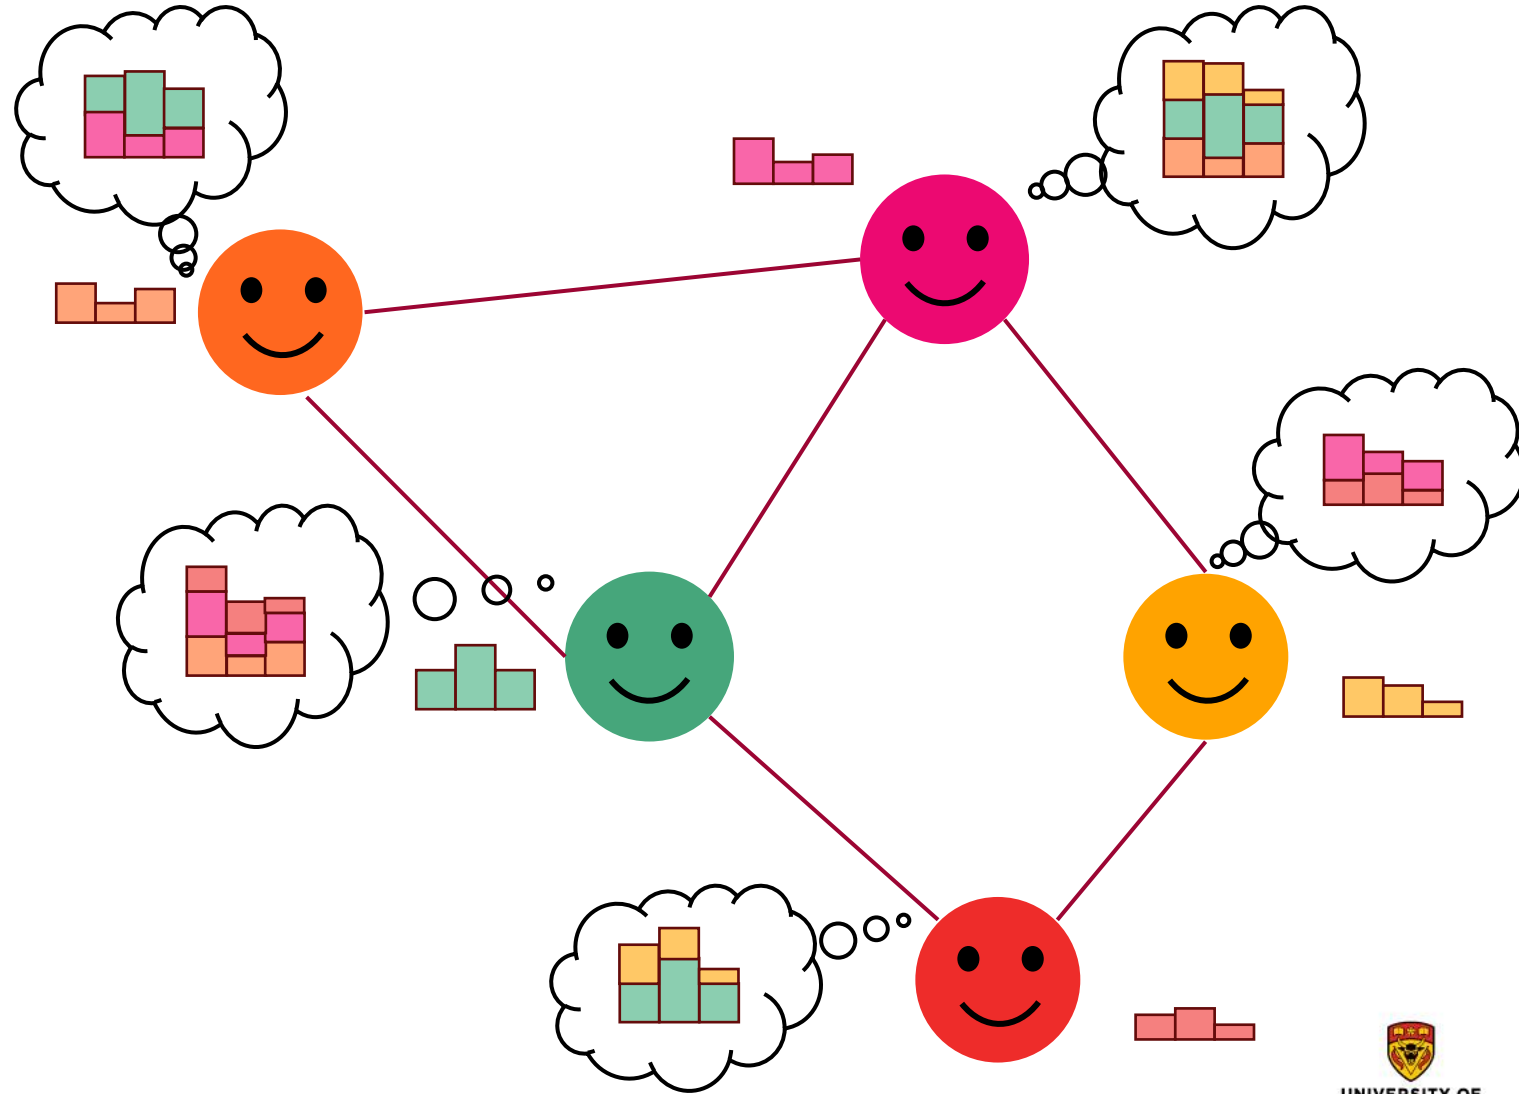
# Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!



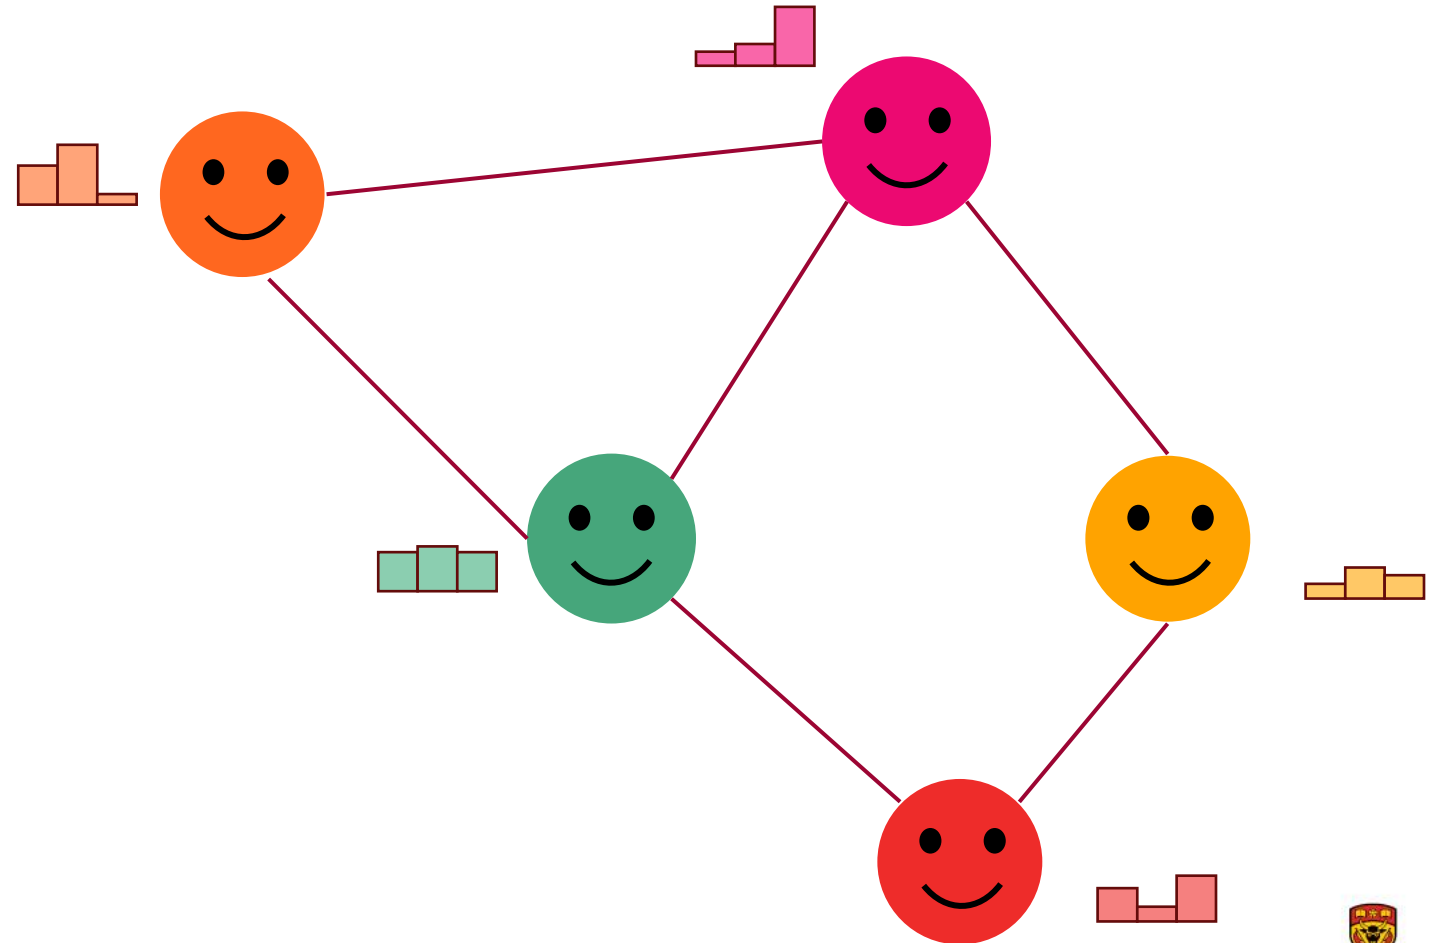
# Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!



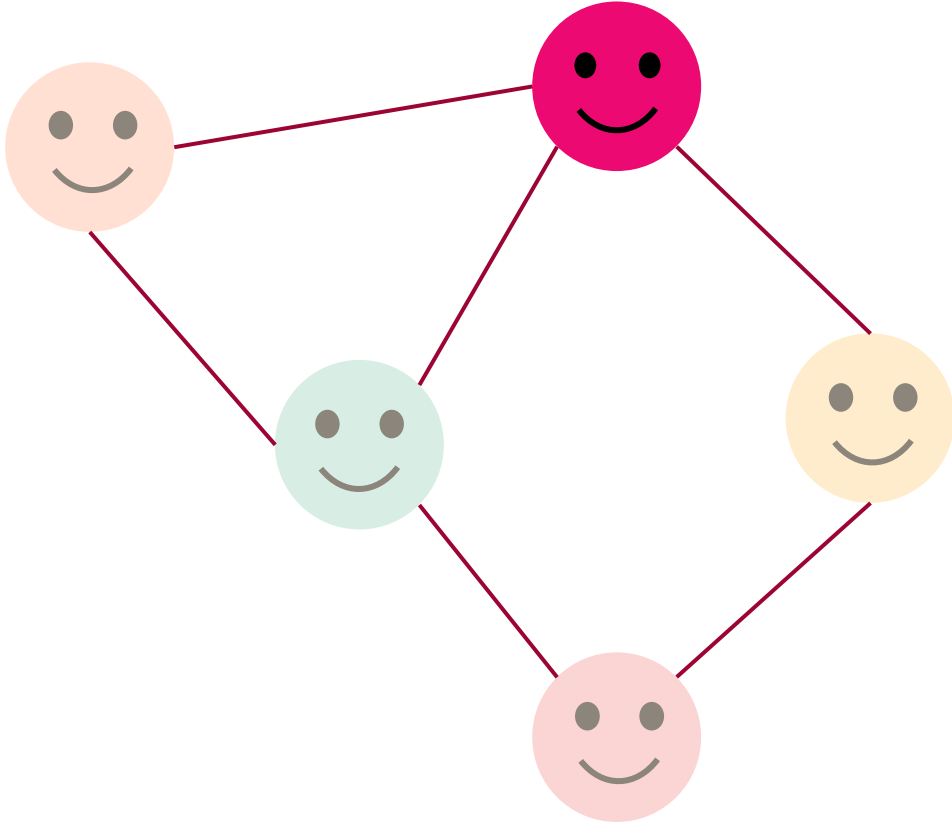
# Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!



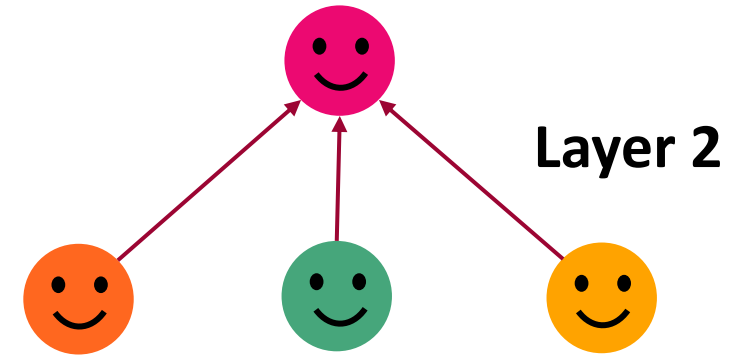
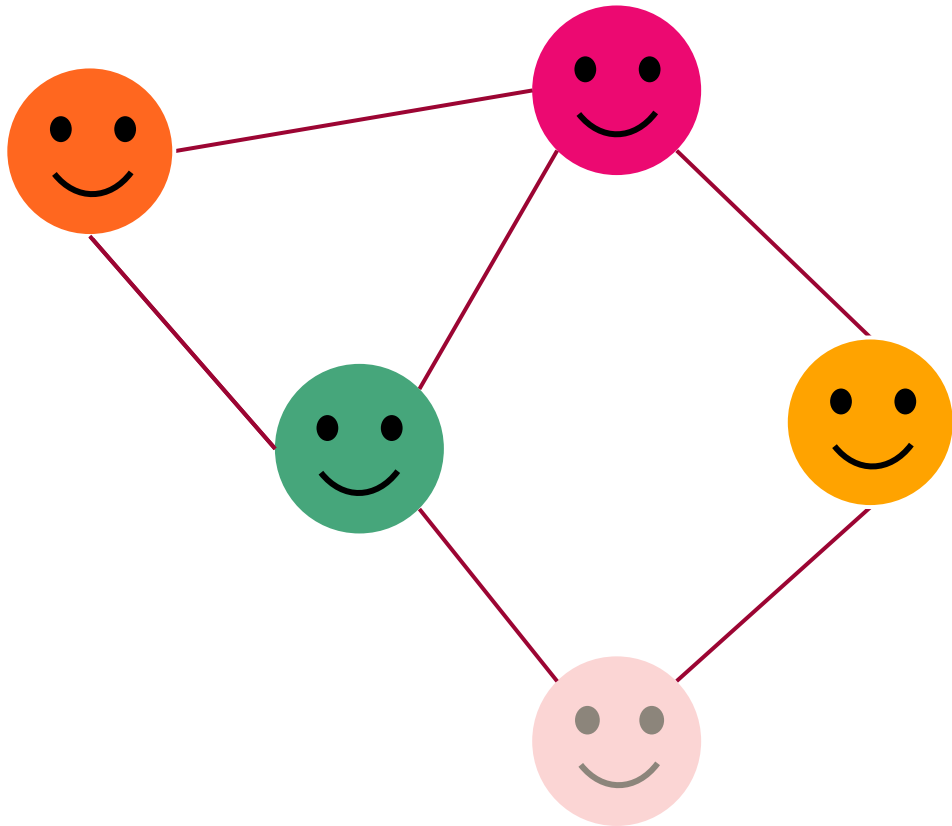
# k-hop neighborhood

- The number of layers we have corresponds to the 'k-hop' neighborhood



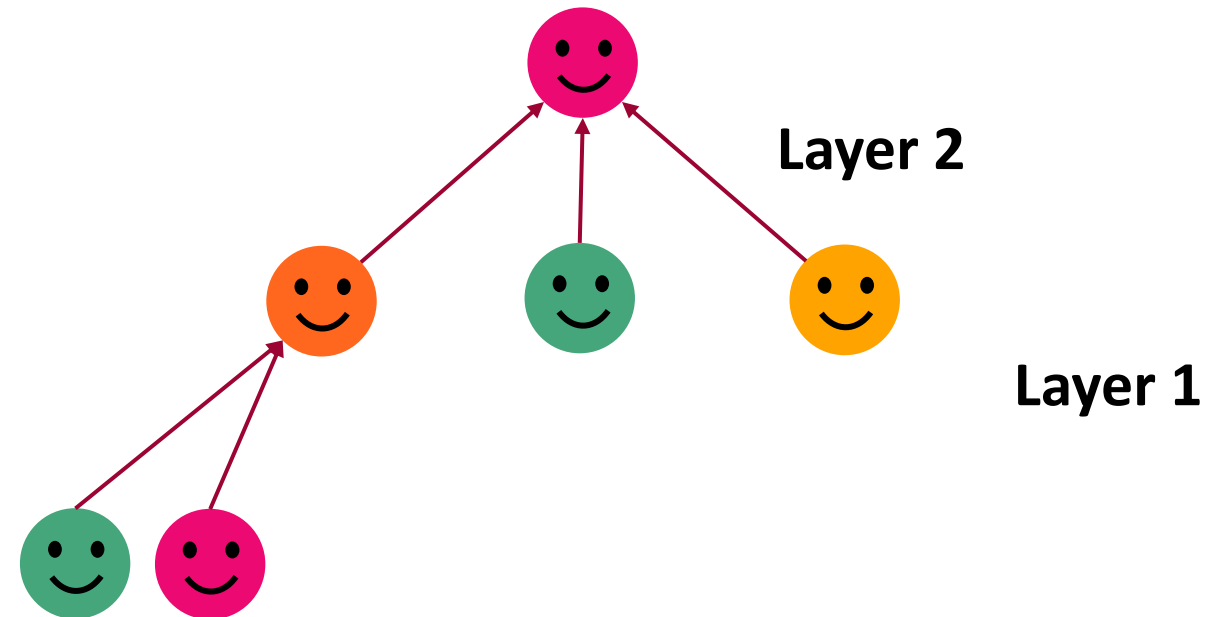
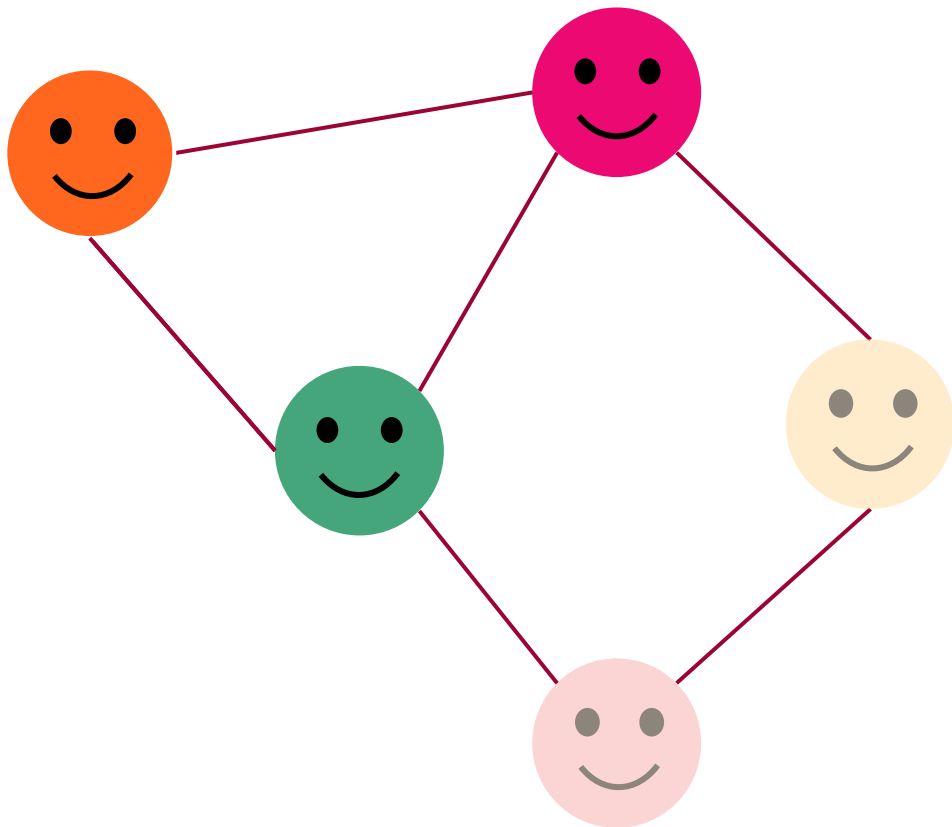
# k-hop neighborhood

- The number of layers we have corresponds to the 'k-hop' neighborhood



# k-hop neighborhood

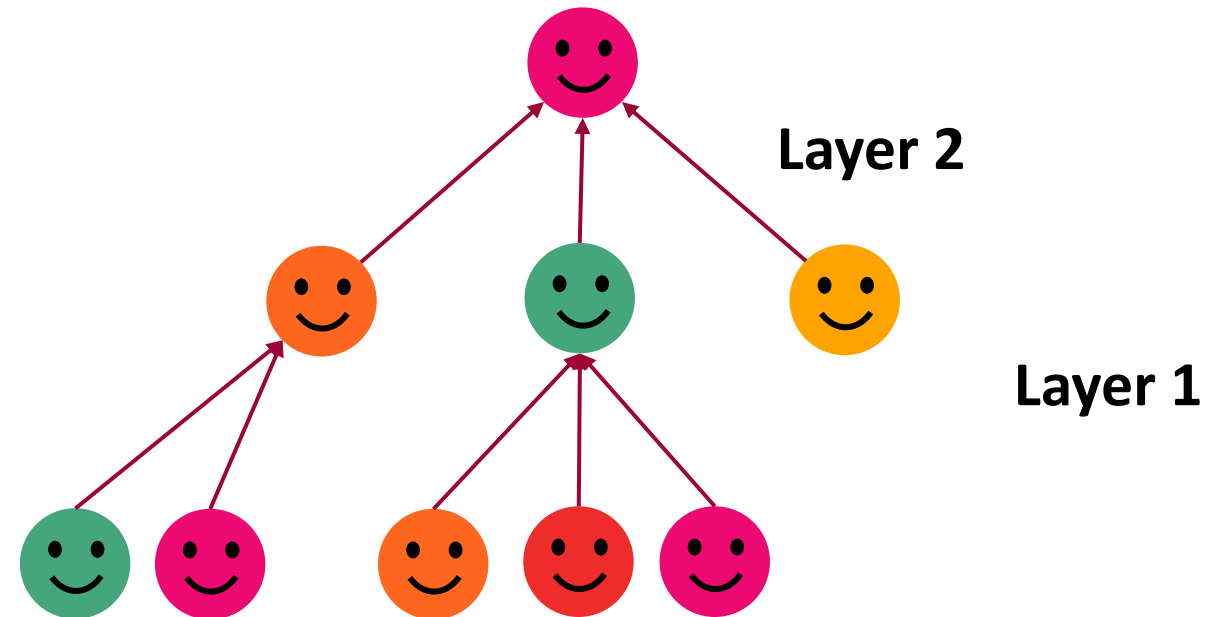
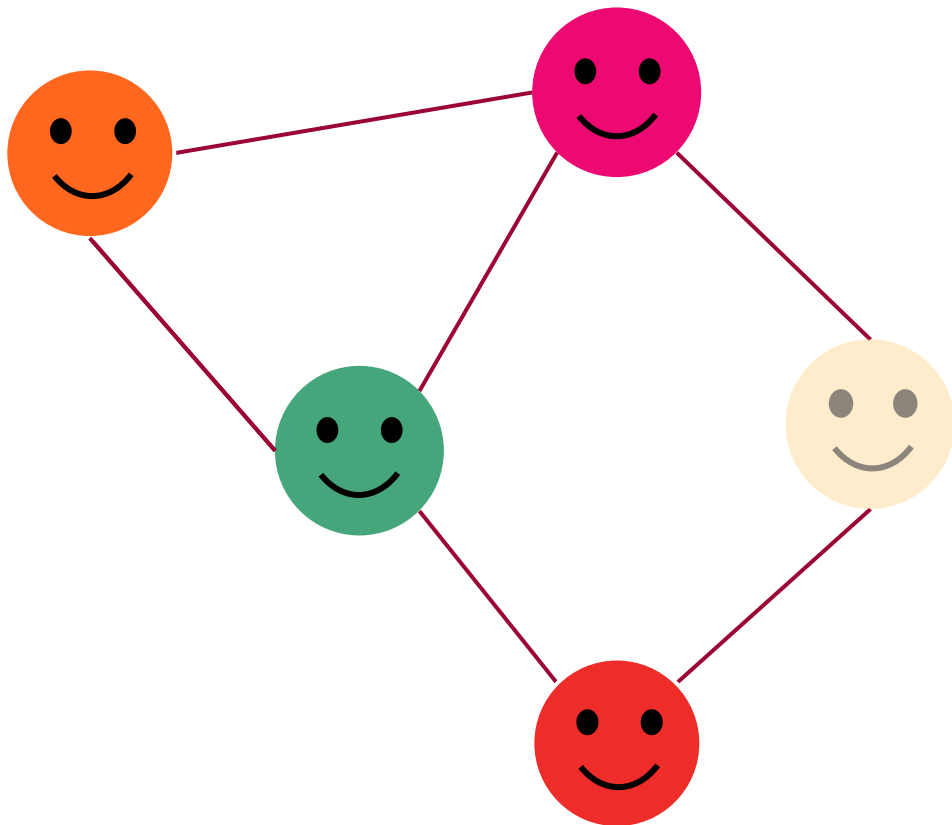
- The number of layers we have corresponds to the 'k-hop' neighborhood





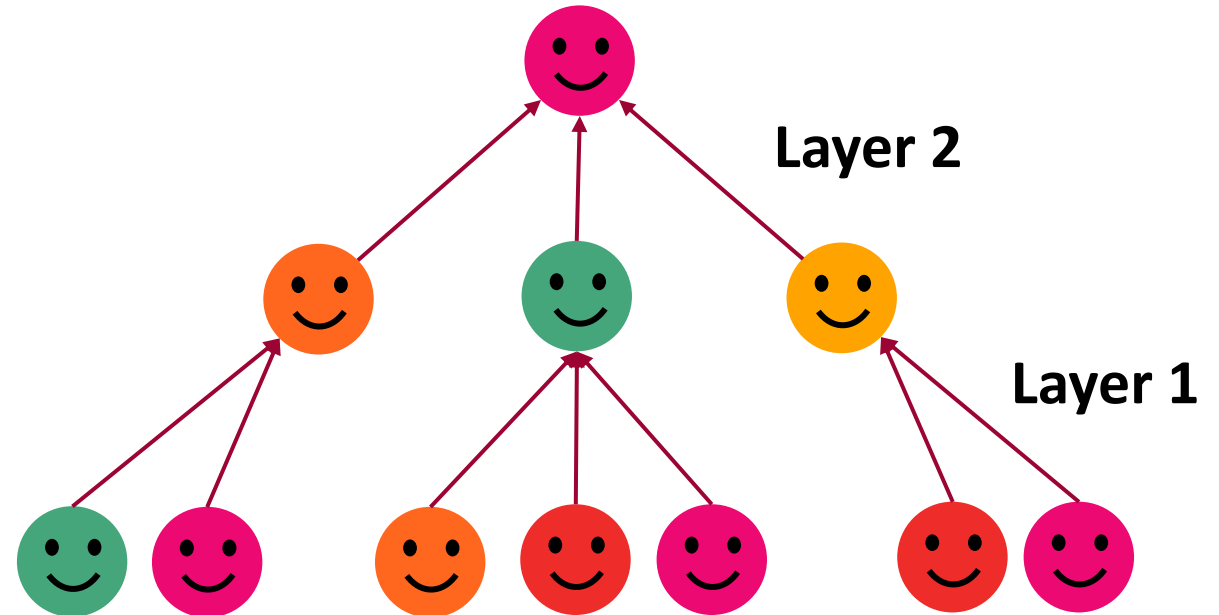
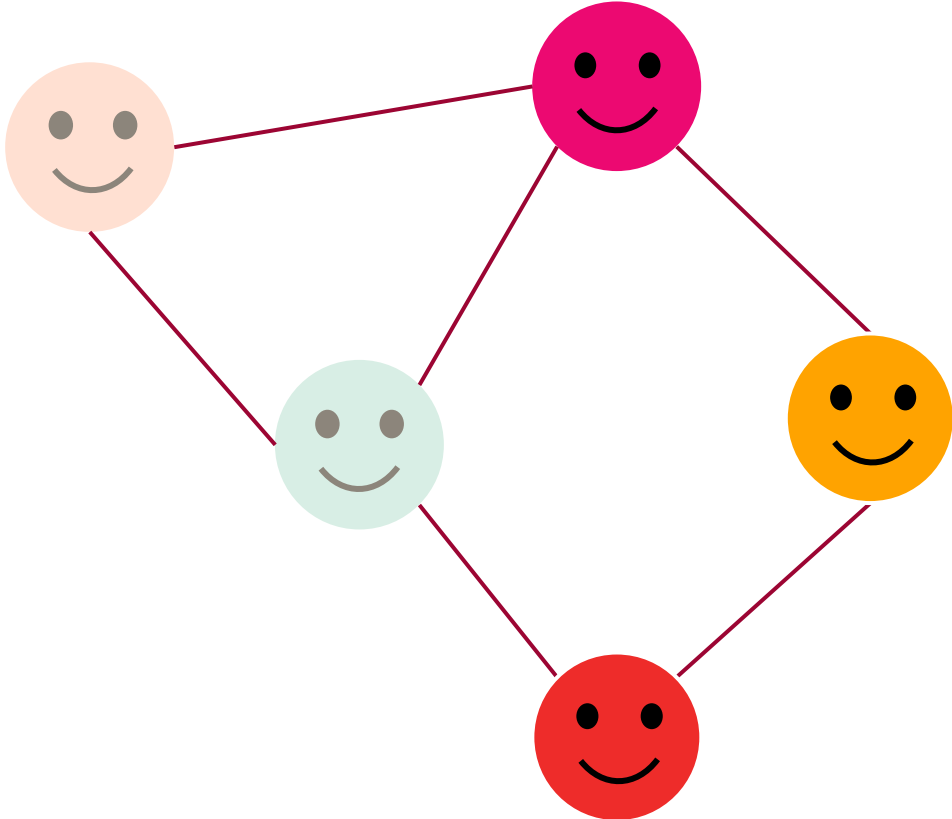
# k-hop neighborhood

- The number of layers we have corresponds to the 'k-hop' neighborhood



# k-hop neighborhood

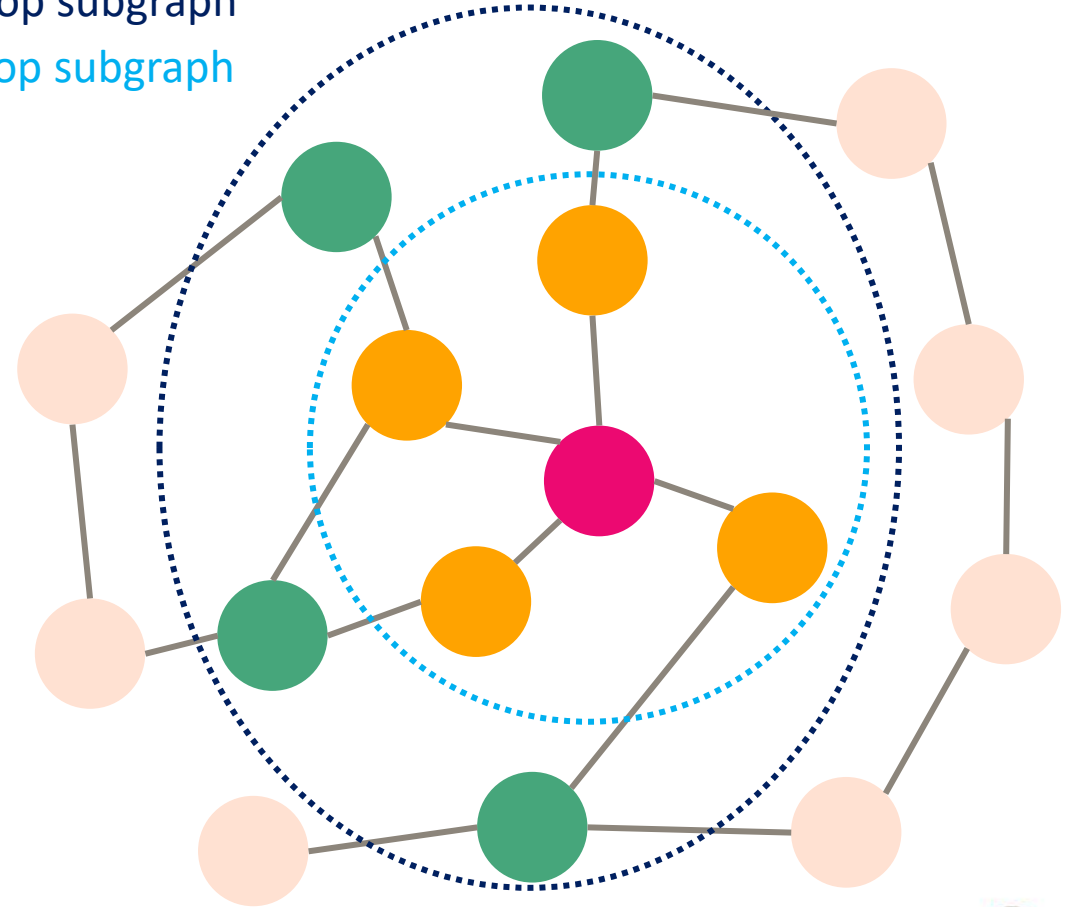
- The number of layers we have corresponds to the 'k-hop' neighborhood



# k-hop neighborhood

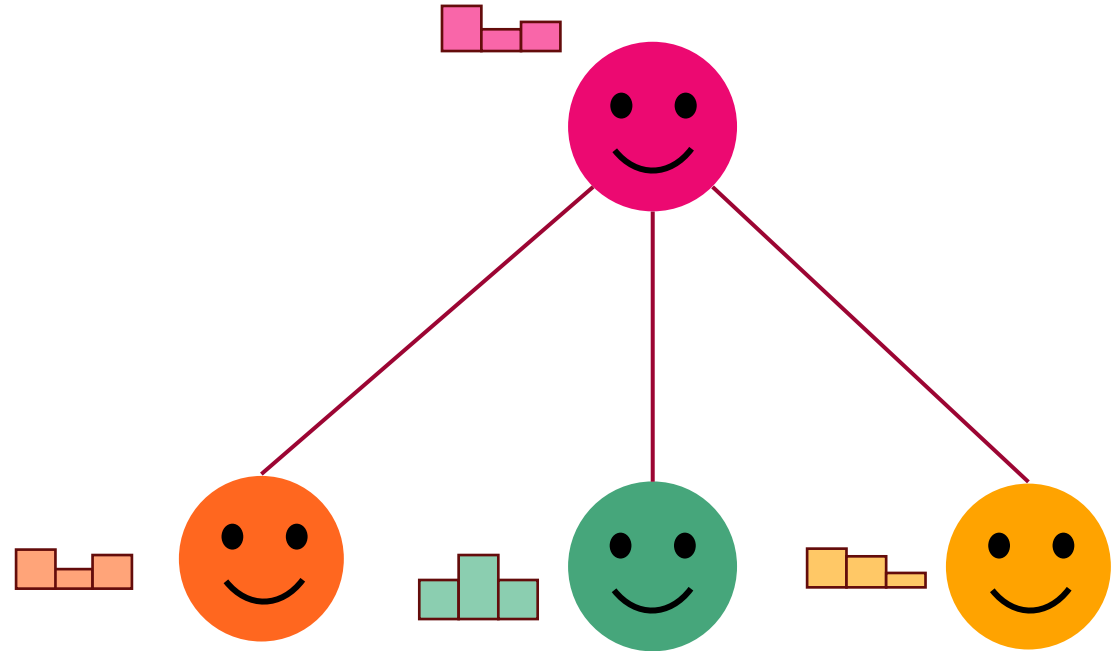
- The more hops, the more distant information each node gets
- More hops is good for global context
- ...but too many hops will wash out local structural information

2-hop subgraph  
1-hop subgraph



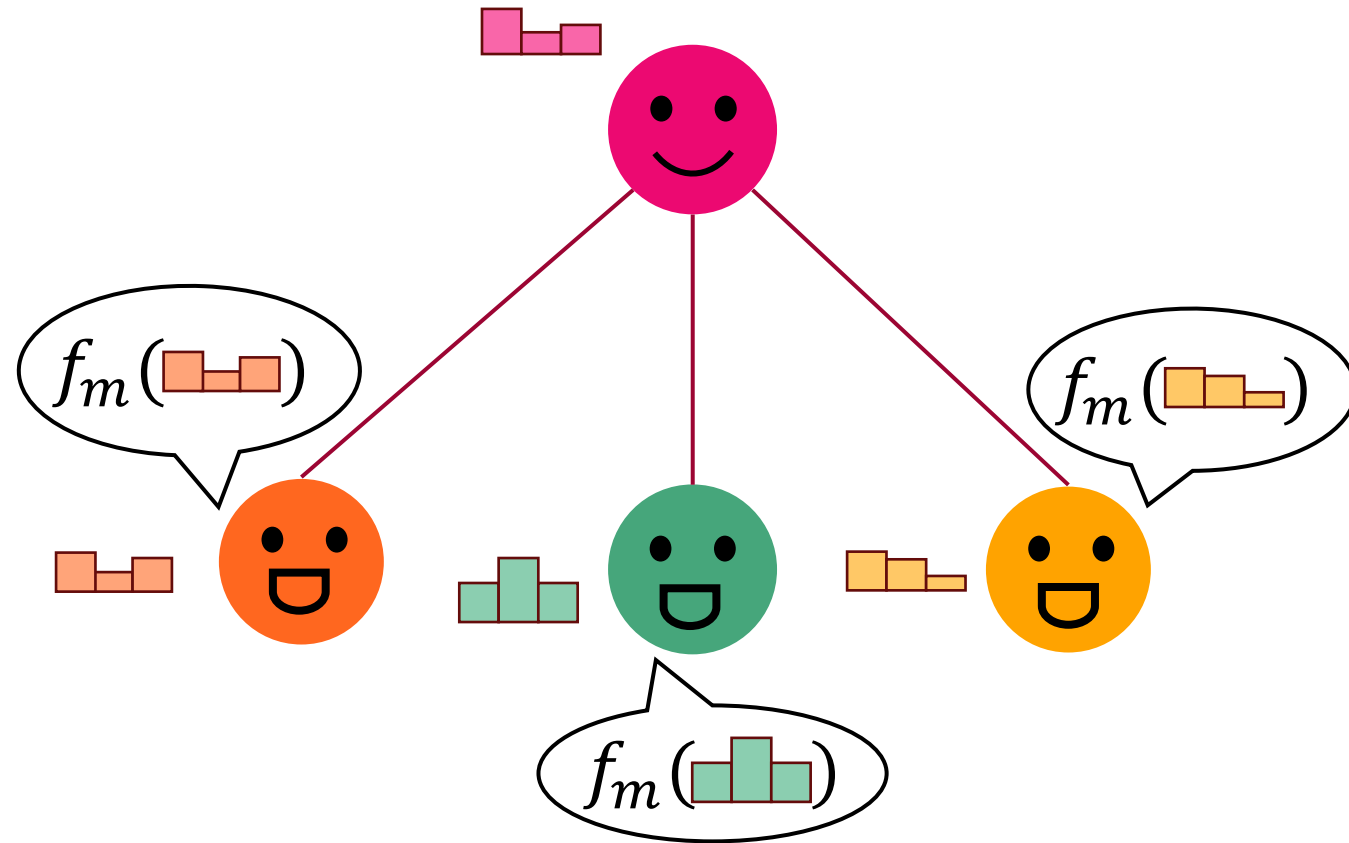
# Trainable parameters

- Can think of it as two places we can have trainable parameters:



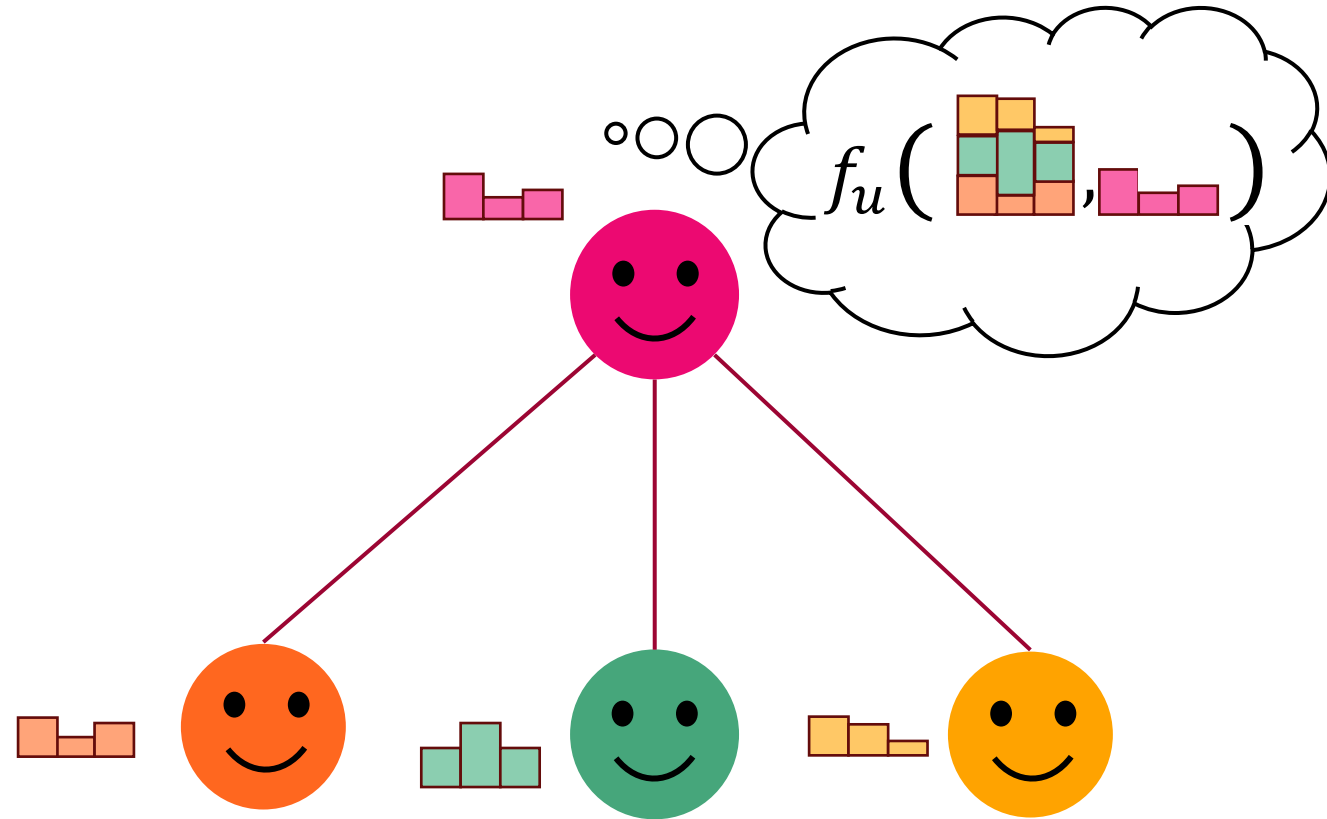
# Trainable parameters

- Can think of it as two places we can have trainable parameters:
  - When passing a message



# Trainable parameters

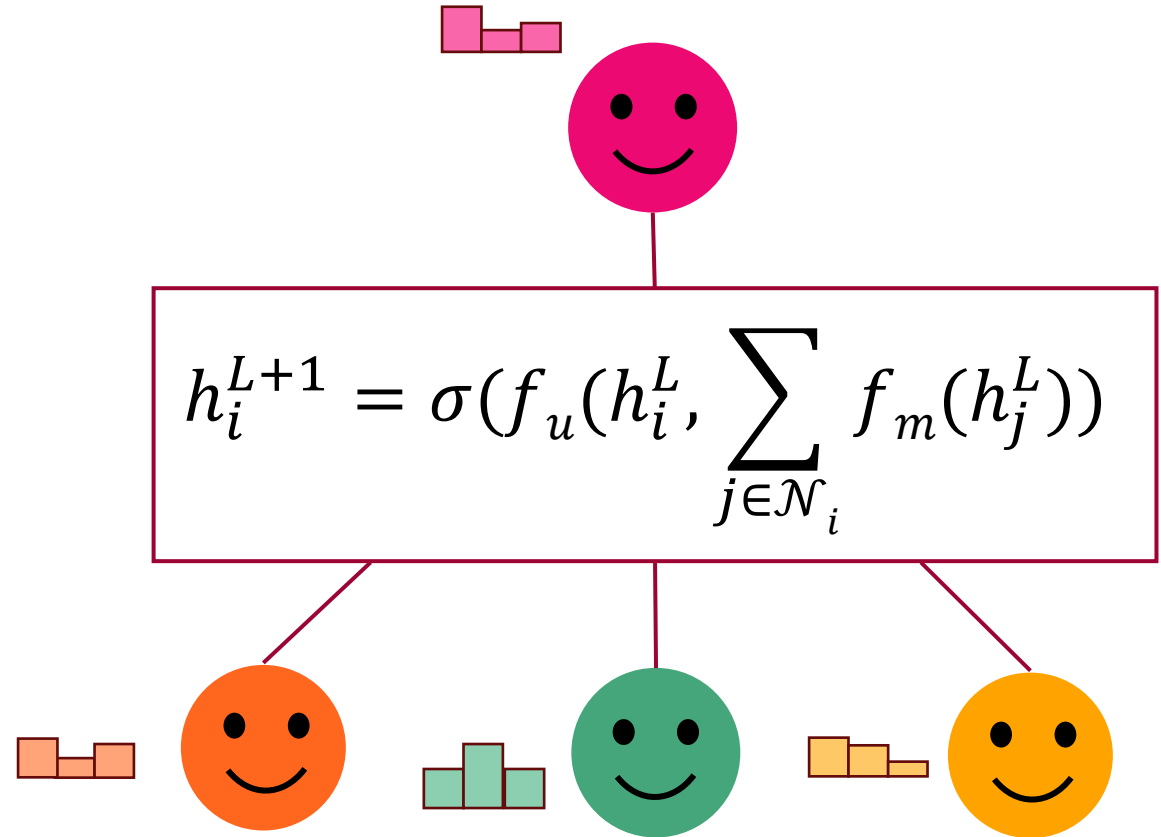
- Can think of it as two places we can have trainable parameters:
  - When passing a message
  - When updating an embedding



# Trainable parameters

- Can think of it as two places we can have trainable parameters:
  - When passing a message
  - When updating an embedding

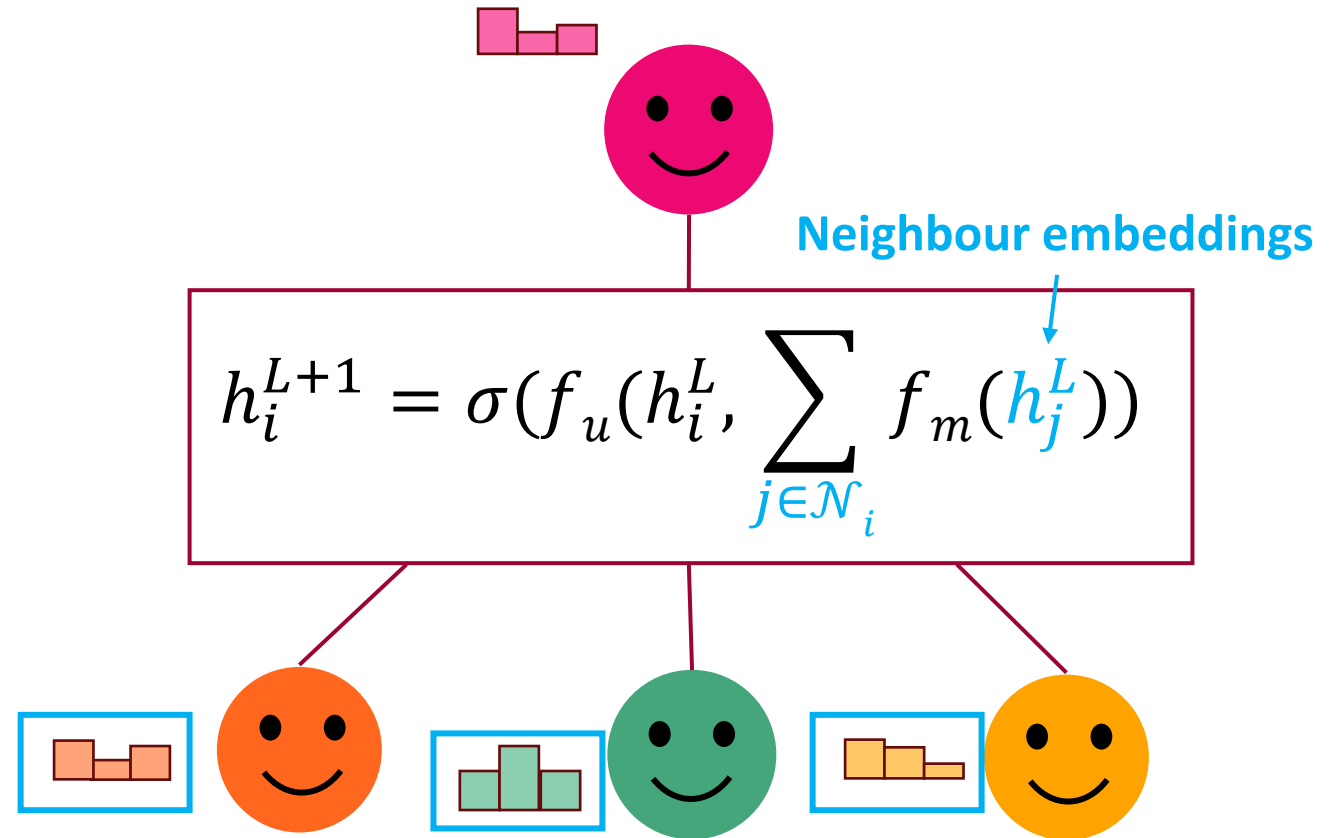
\* Note: The trainable parameters are shared across ALL nodes!



# Trainable parameters

- Can think of it as two places we can have trainable parameters:
  - When passing a message
  - When updating an embedding

\* Note: The trainable parameters are shared across ALL nodes!

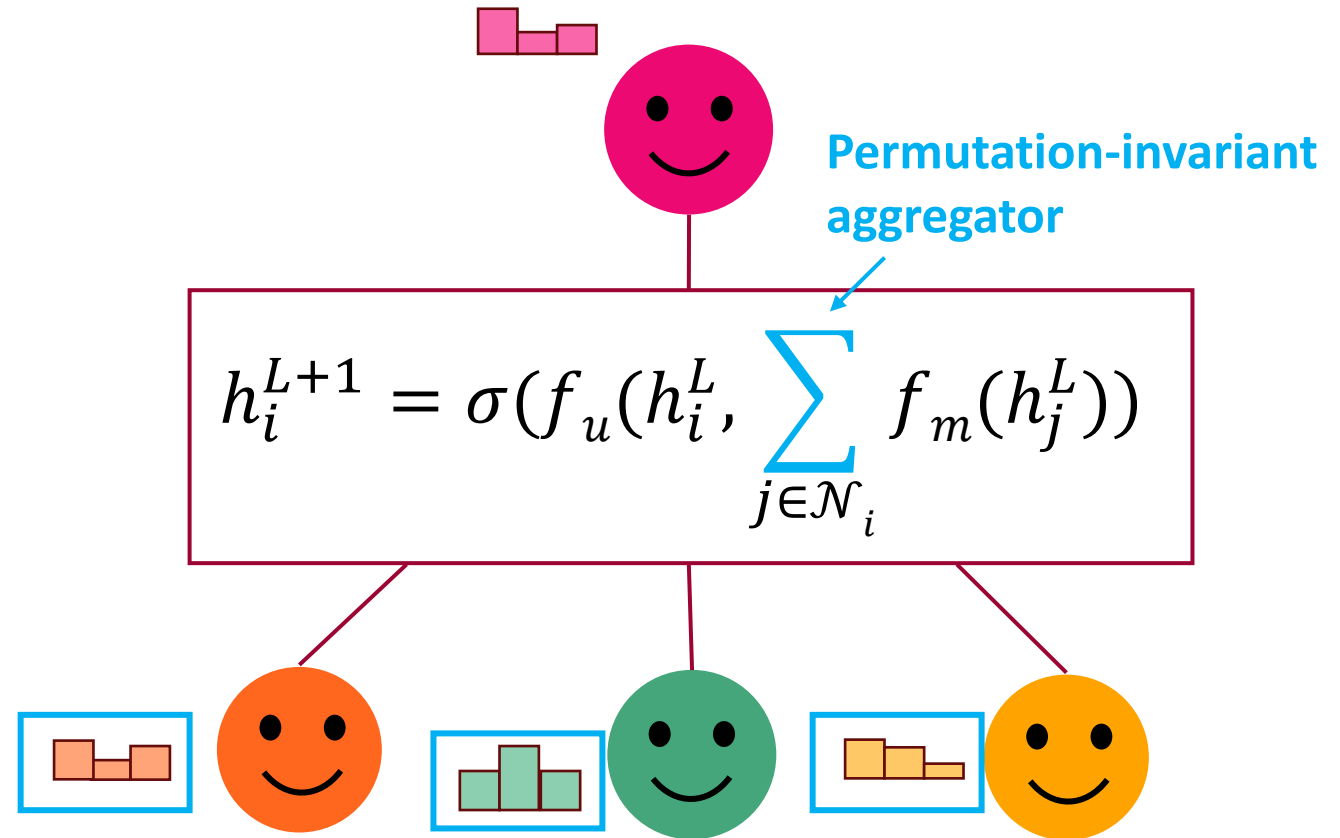




# Trainable parameters

- Can think of it as two places we can have trainable parameters:
  - When passing a message
  - When updating an embedding

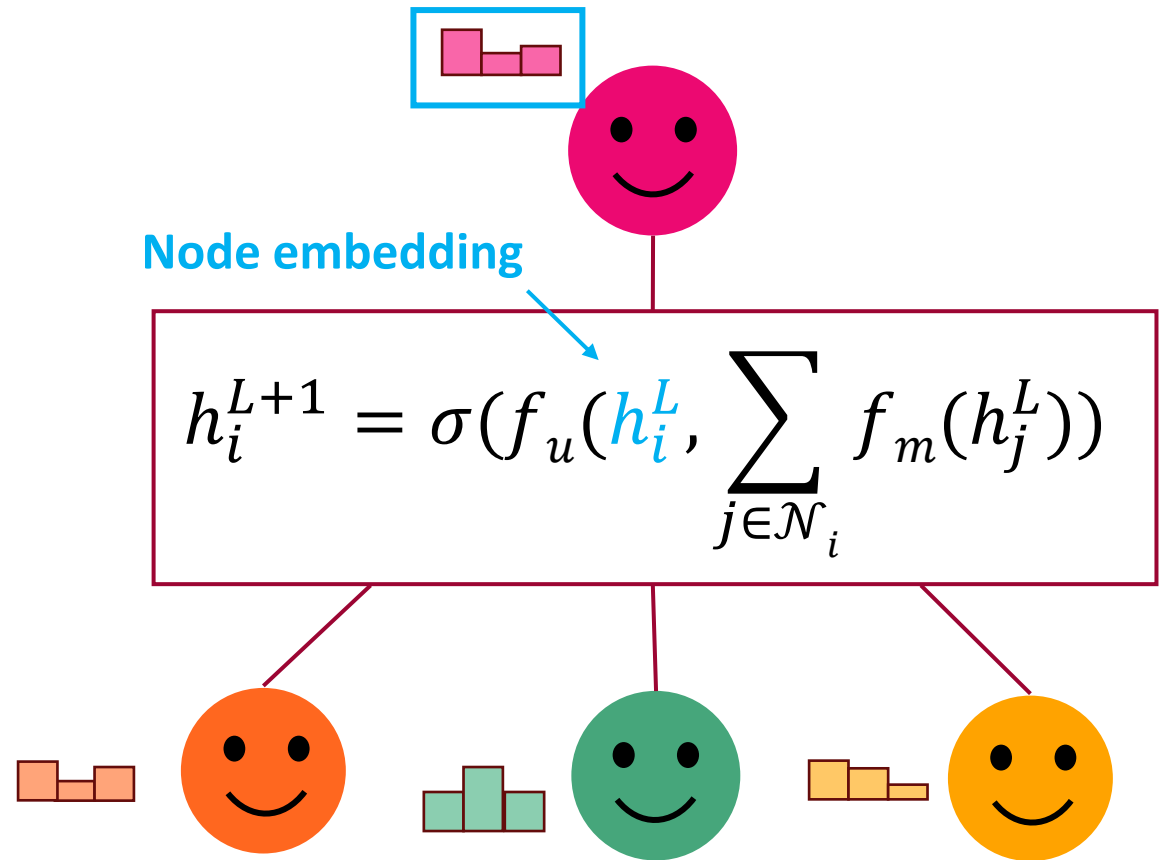
\* Note: The trainable parameters are shared across ALL nodes!



# Trainable parameters

- Can think of it as two places we can have trainable parameters:
  - When passing a message
  - When updating an embedding

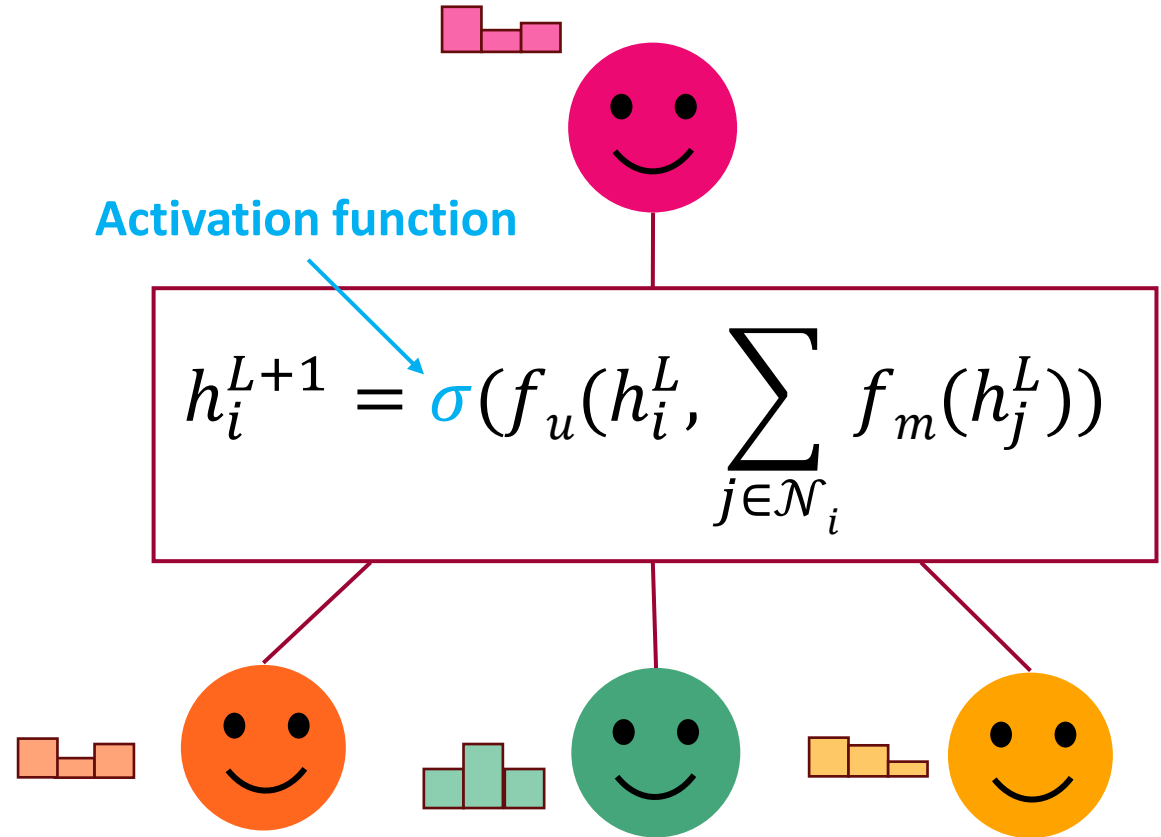
\* Note: The trainable parameters are shared across ALL nodes!



# Trainable parameters

- Can think of it as two places we can have trainable parameters:
  - When passing a message
  - When updating an embedding

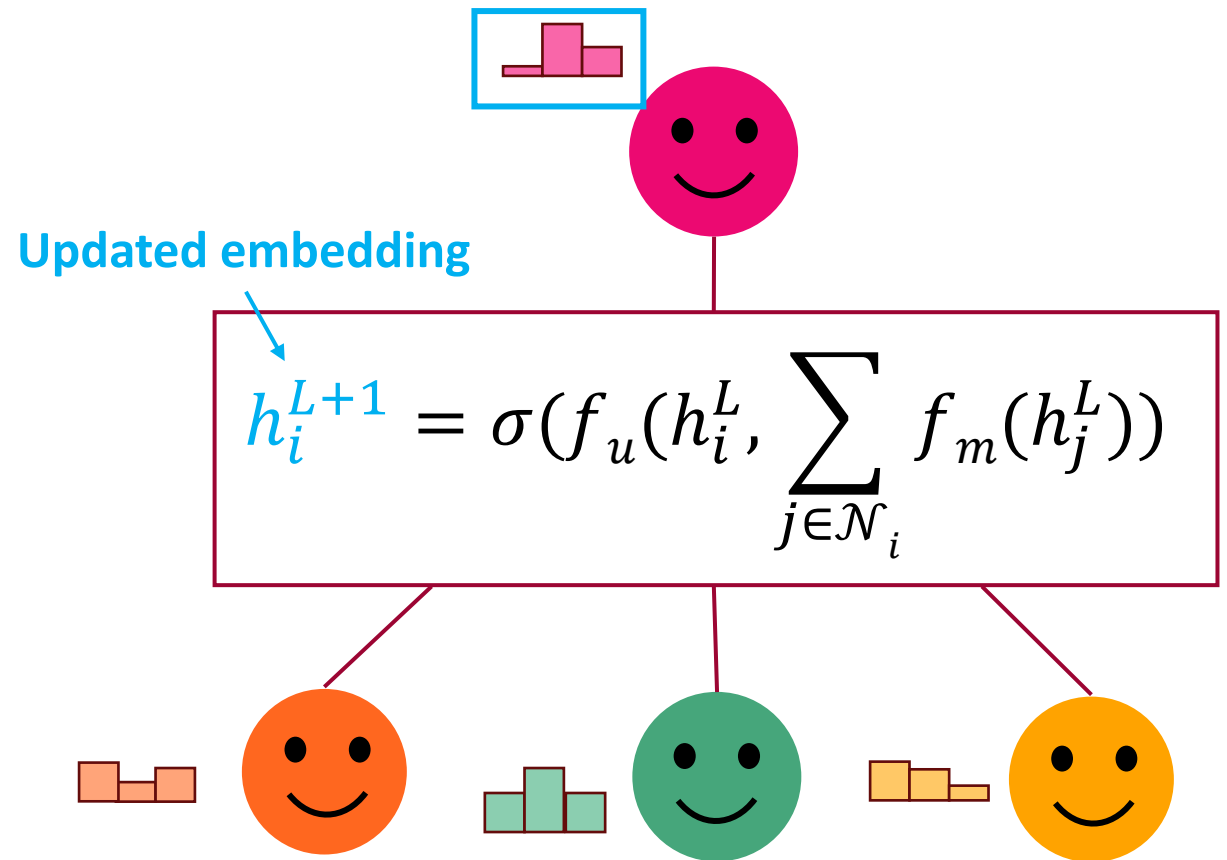
\* Note: The trainable parameters are shared across ALL nodes!



# Trainable parameters

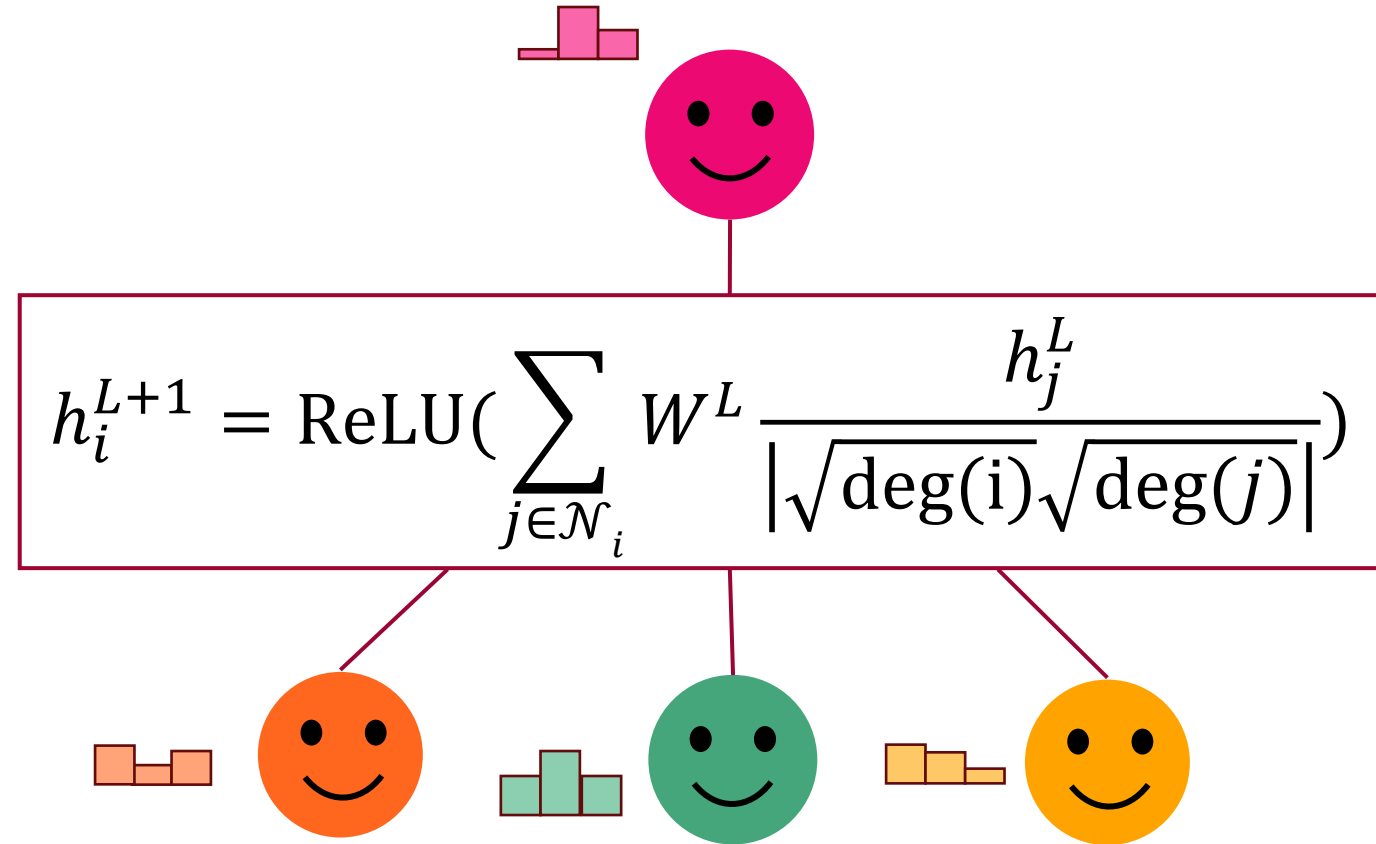
- Can think of it as two places we can have trainable parameters:
  - When passing a message
  - When updating an embedding

\* Note: The trainable parameters are shared across ALL nodes!



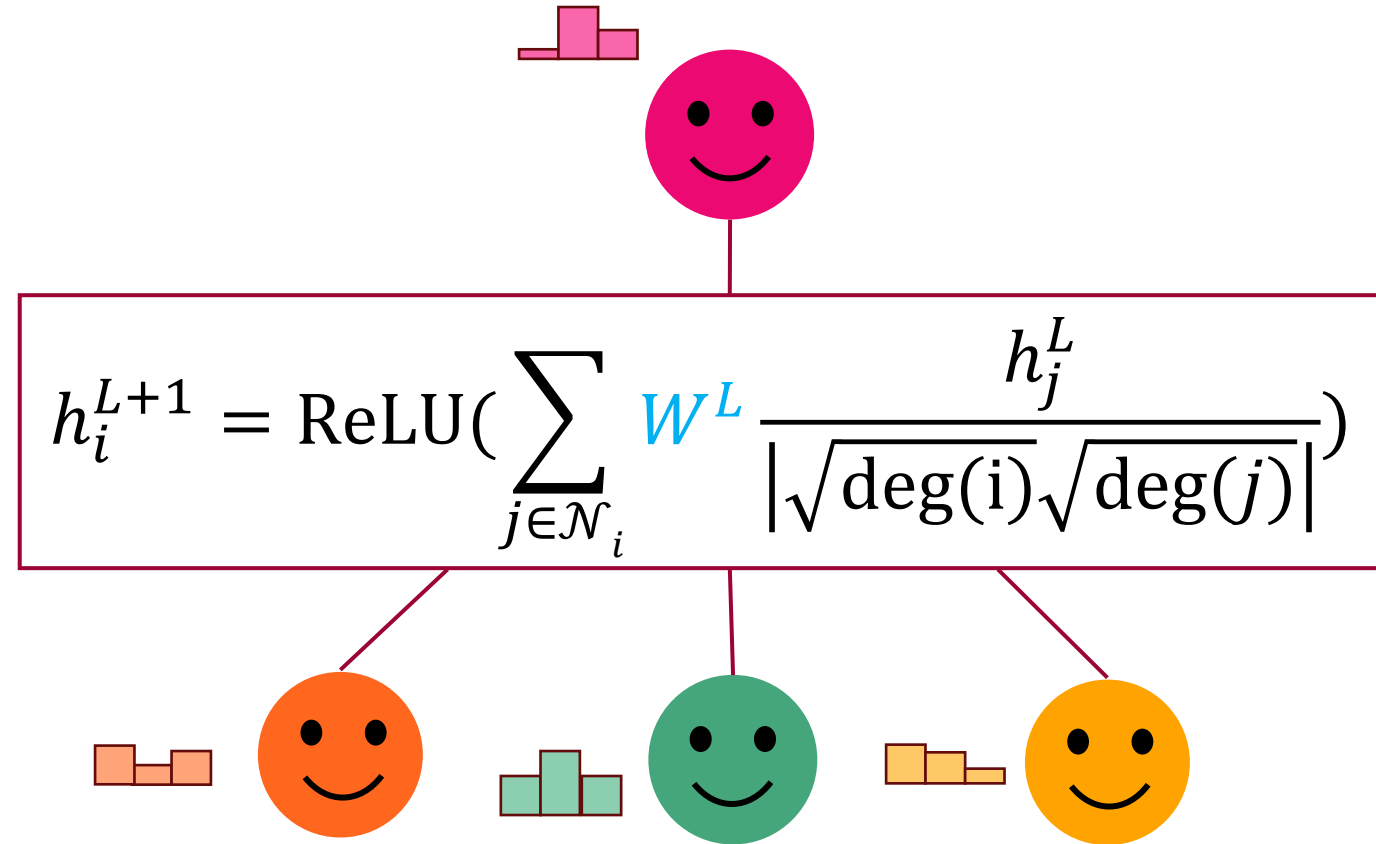
# GNN layer example: Graph Convolutional Network

- Graph convolutional network (GCN) learns a simple matrix of weights
  - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel



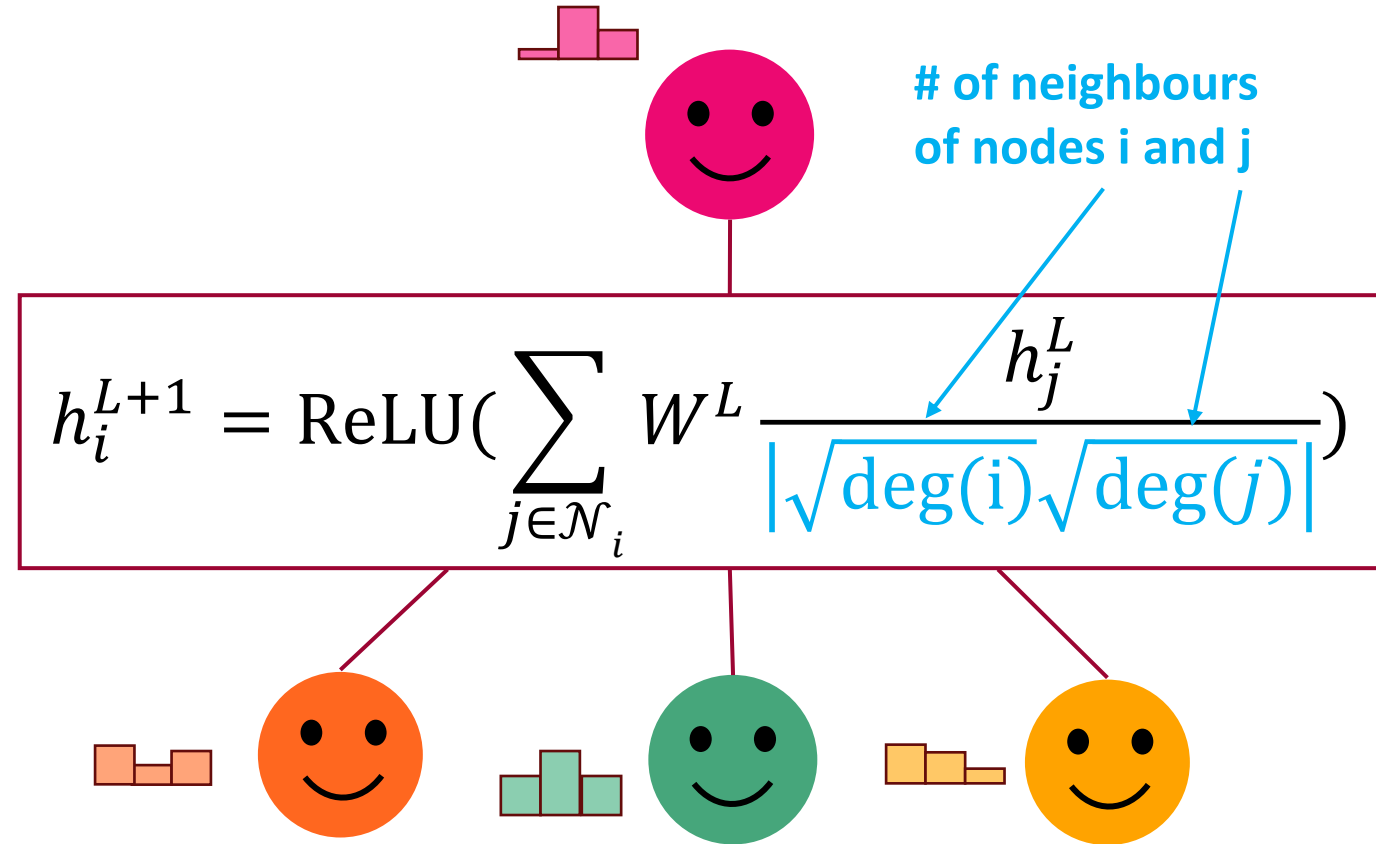
# GNN layer example: Graph Convolutional Network

- Graph convolutional network (GCN) learns a simple matrix of weights
  - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel



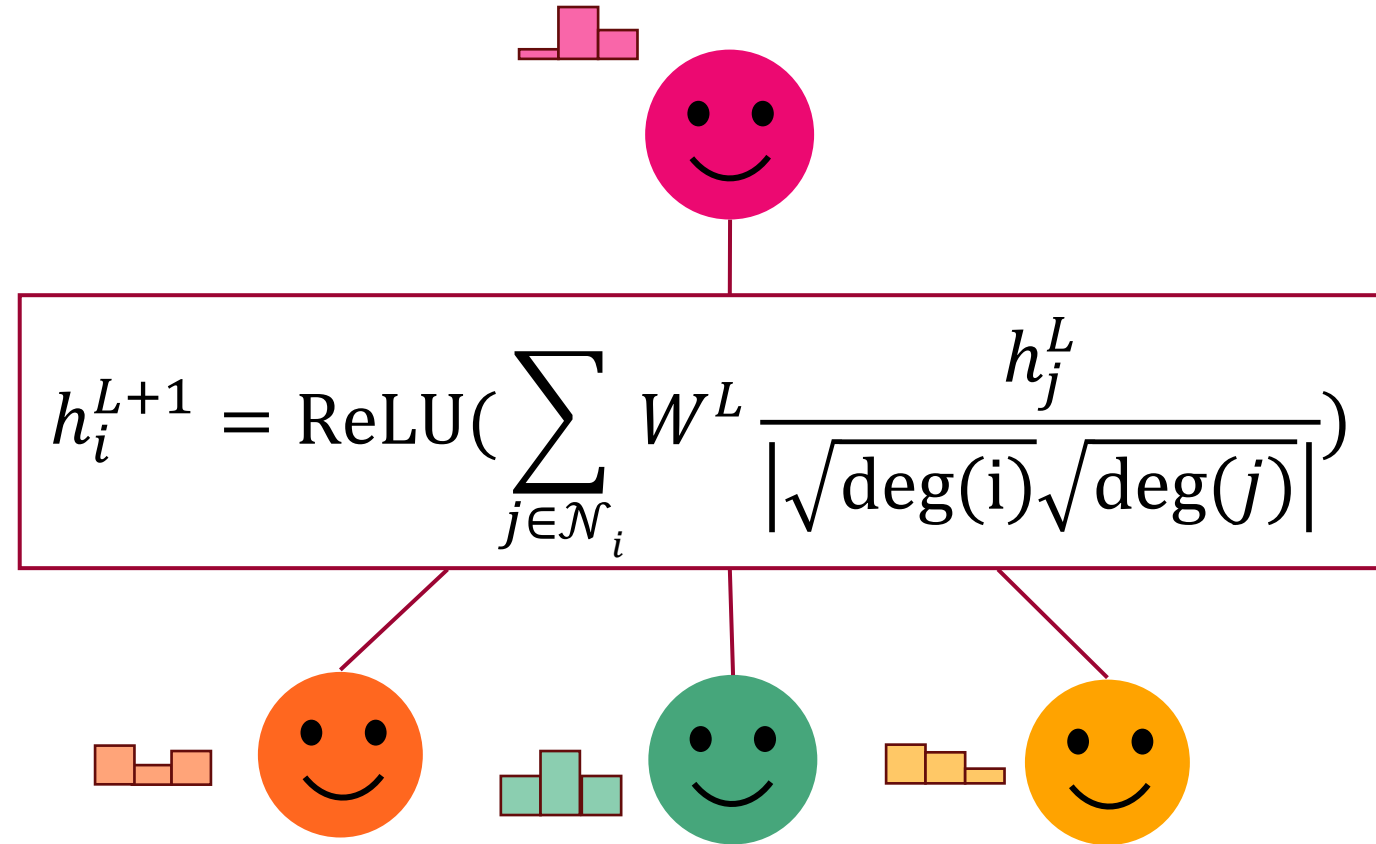
# GNN layer example: Graph Convolutional Network

- Graph convolutional network (GCN) learns a simple matrix of weights
  - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel



# GNN layer example: Graph Convolutional Network

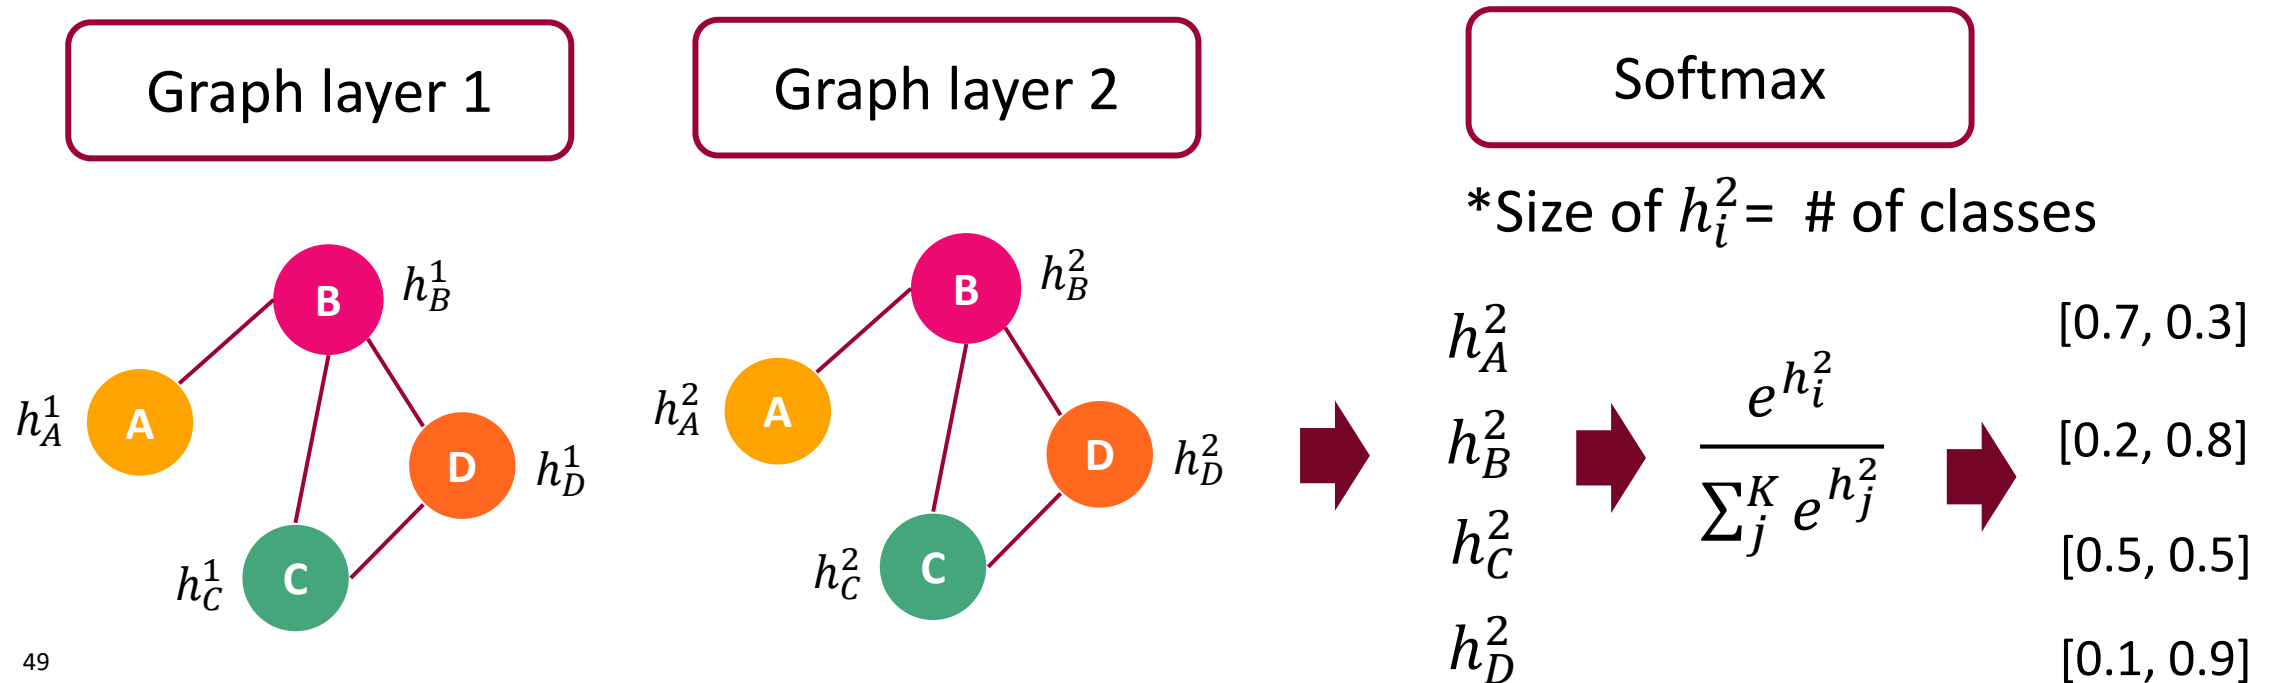
- Graph convolutional network (GCN) learns a simple matrix of weights
  - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel
- GCN assumes self-edges!
  - So, each node is its own neighbor 🤔





# Putting it all together

- We *could* just learn node/edge/graph embeddings e.g. node2vec
- To tailor network for a specific task, we add an appropriate head
- E.g. Node classification:



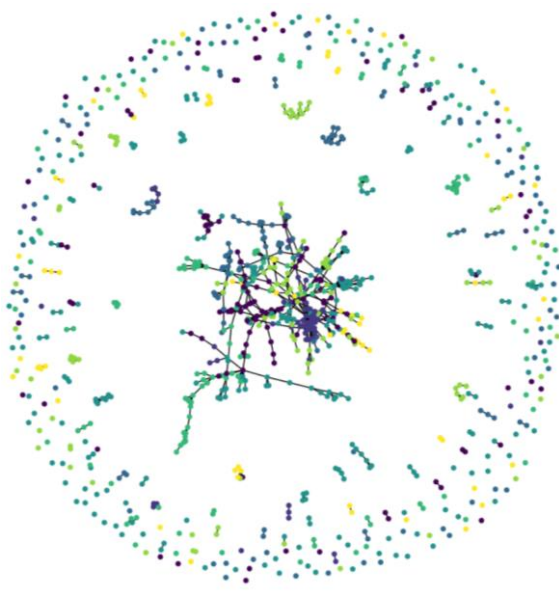
# Summary

- Graph structured data is one of the most general types of data that encompasses a broad range of applications
- Graph neural networks can be used to solve a variety of tasks
- Since graph data is inherently different from other types of data, we need a unique framework to handle it
- Message passing allows for all the information in a graph to be effectively harnessed

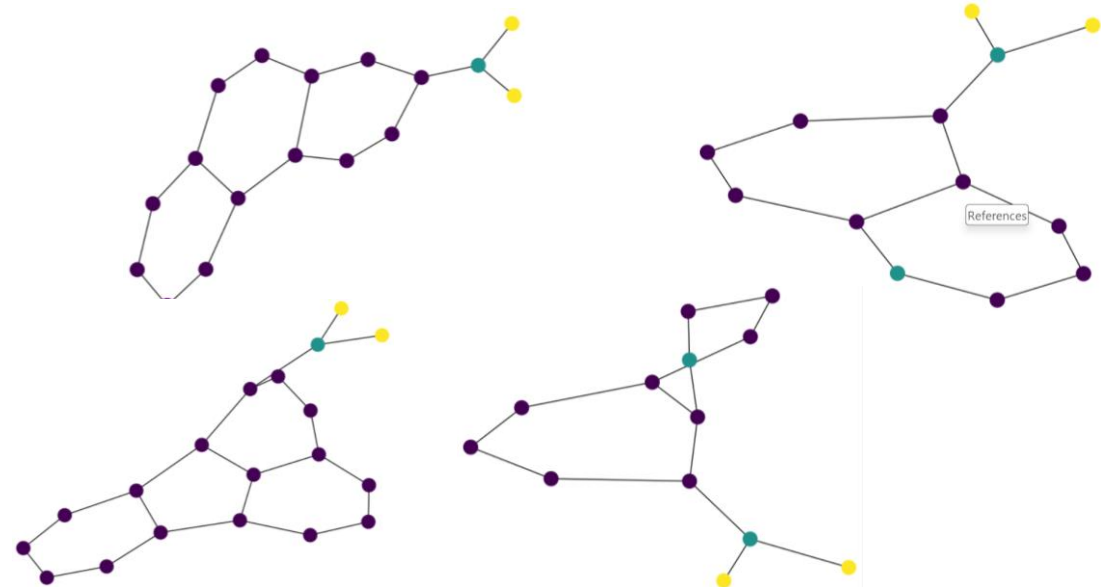
# Next class:

- We will learn how to work with graphs and build models using pytorch-geometric!

**Cora dataset (Citations)**



**TU MUTAG Dataset (Molecules)**



# References

- J. Leskovec. [CS224W: Machine Learning with Graphs](#). \*\* Great option if you want a whole COURSE on graph-based machine learning
- W. L. Hamilton. Graph Representation Learning. *Synthesis Lectures on Artificial Intelligence and Machine Learning*, 14(3), pages 1-159. 2020.
- Sanchez-Lengeling, et al., "A Gentle Introduction to Graph Neural Networks", Distill, 2021.
- T. N. Kipf et al., Semi-Supervised Classification with graph Convolutional Networks, ICLR 2017.
- K. Xu et al., How Powerful are Graph Neural Networks? ICLR, 2019.
- W. L. Hamilton et al., Inductive representation Learning on Large Graphs. NIPS, 2017.
- J. Zhou et al., Graph neural networks: A review of methods and applications. AI Open 1. Pages 57-81. 2020.
- Daigavane, et al., "Understanding Convolutions on Graphs", Distill, 2021.
- P. Lippe. [Tutorial 7: Graph Neural Networks — UvA DL Notebooks v1.2 documentation \(uvadlc-notebooks.readthedocs.io\)](#). 2022.
- M. N. Bernstein. [Graph convolutional neural networks - Matthew N. Bernstein \(mbernste.github.io\)](#). 2023.
- Z. Wu et al., A Comprehensive Survey on Graph Neural Networks. IEEE Trans. On Neural networks and Learning Systems. 32(1). 2021.
- R. Anand. [Math Behind Graph Neural Networks - Rishabh Anand \(rish-16.github.io\)](#). 2022.
- D. Grattarola. [A practical introduction to GNNs - Part 2 – Daniele Grattarola](#). 2021.
- E. Benjaminson. [Understanding Message Passing in GNNs – Emma Benjaminson – Data Scientist \(sassafra13.github.io\)](#). 2022.
- T. Masui. [Graph Neural Networks with PyG on Node Classification, Link Prediction, and Anomaly Detection | by Tomonori Masui | Towards Data Science](#). 2022.

# Thank you!

